

Main Memory (Part II)

Amir H. Payberah
amir@sics.se

Amirkabir University of Technology
(Tehran Polytechnic)




Reminder


Reminder (1/3)

- ▶ External fragmentation vs. internal fragmentation
- ▶ **Compaction**: shuffle memory contents to place all free memory together in one large block.
- ▶ Other solutions:
 - Segmentation
 - Paging

Reminder (2/3)

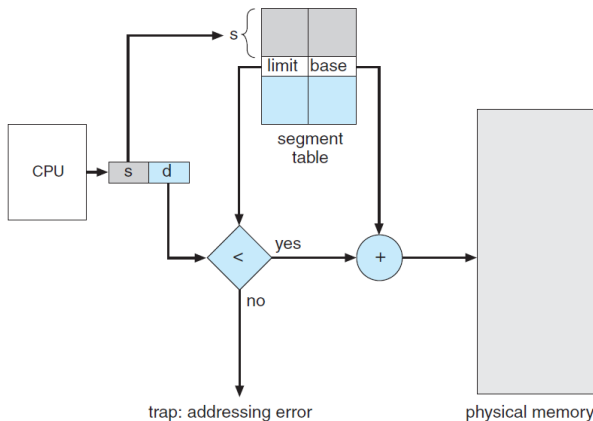


user space



physical memory space

Reminder (3/3)



Paging

Paging vs. Segmentation

- ▶ **Segmentation** and **paging**, both, permit the physical address space of a process to be **noncontiguous**.

Paging vs. Segmentation

- ▶ **Segmentation** and **paging**, both, permit the **physical address space** of a process to be **noncontiguous**.
- ▶ **Paging** avoids external fragmentation and the need for compaction, whereas segmentation does not.

Paging (1/2)

- ▶ **Physical address space** of a process can be **noncontiguous**; process is allocated physical memory whenever the latter is available.
 - Avoids **external fragmentation**
 - Avoids problem of **varying sized** memory chunks

Paging (1/2)

- ▶ **Physical address space** of a process can be **noncontiguous**; process is allocated physical memory whenever the latter is available.
 - Avoids **external fragmentation**
 - Avoids problem of **varying sized** memory chunks
- ▶ Divide **physical memory** into fixed-sized blocks called **frames**.
 - Size is **power of 2**, between 512 bytes and 16 Mbytes.

Paging (1/2)

- ▶ Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available.
 - Avoids external fragmentation
 - Avoids problem of varying sized memory chunks
- ▶ Divide physical memory into fixed-sized blocks called frames.
 - Size is power of 2, between 512 bytes and 16 Mbytes.
- ▶ Divide logical memory into blocks of same size called pages.

Paging (2/2)

- ▶ Keep track of all free frames.

Paging (2/2)

- ▶ Keep track of all free frames.
- ▶ To run a program of size N pages, need to find N free frames and load program.

Paging (2/2)

- ▶ Keep track of all free frames.
- ▶ To run a program of size N pages, need to find N free frames and load program.
- ▶ Set up a page table to translate logical to physical addresses.

Paging (2/2)

- ▶ Keep track of all free frames.
- ▶ To run a program of size N pages, need to find N free frames and load program.
- ▶ Set up a page table to translate logical to physical addresses.
- ▶ Still have internal fragmentation.

Address Translation Scheme


- ▶ Address generated by CPU (logical address) is divided into **two parts**:

Address Translation Scheme

- ▶ Address generated by CPU (logical address) is divided into two parts:
- ▶ Page number (p): used as an index into a page table that contains base address of each page in physical memory.


Address Translation Scheme

- ▶ Address generated by CPU (logical address) is divided into **two parts**:
- ▶ **Page number** (p): used as an **index into a page table** that contains **base address of each page** in **physical memory**.
- ▶ **Page offset** (d): combined with base address to define the physical memory address that is sent to the memory unit.



Address Translation Scheme

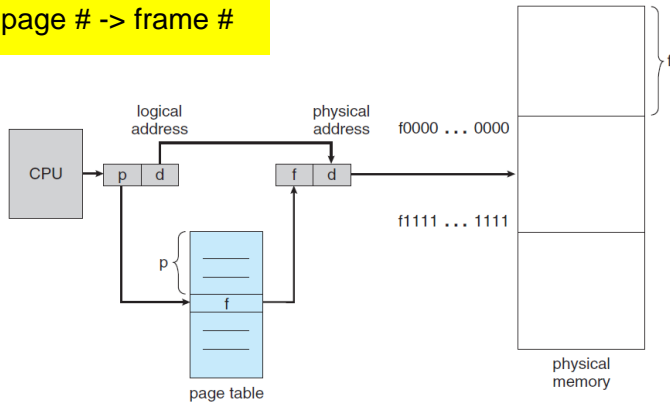
- ▶ Address generated by CPU (logical address) is divided into two parts:
- ▶ Page number (p): used as an index into a page table that contains base address of each page in physical memory.
- ▶ Page offset (d): combined with base address to define the physical memory address that is sent to the memory unit.



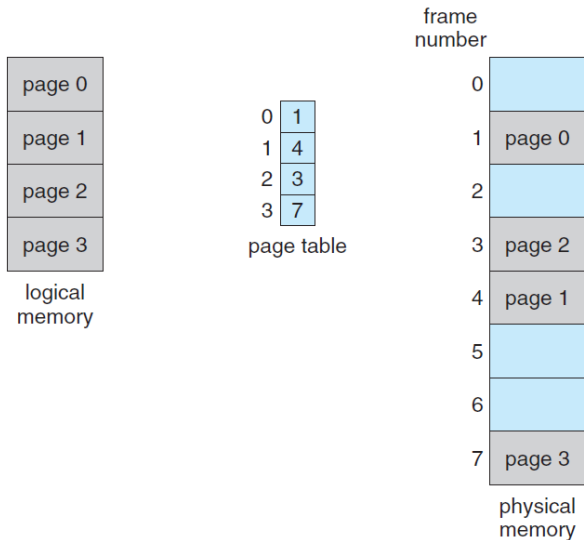
- ▶ For given logical address space 2^m and page size 2^n .

Paging Hardware

page # -> frame #



Paging Model of Logical and Physical Memory



Paging Example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

- $n = 2$ and $m = 4$, 32-byte memory and 4-byte pages

Free Frames



Paging Example - Internal Fragmentation

- ▶ Page size = 2048 bytes
- ▶ Process size = 72766 bytes

Paging Example - Internal Fragmentation

- ▶ Page size = 2048 bytes
- ▶ Process size = 72766 bytes
- ▶ $\frac{72766}{2048} = 35 \text{ pages} + 1086 \text{ bytes}$

Paging Example - Internal Fragmentation

- ▶ Page size = 2048 bytes
- ▶ Process size = 72766 bytes
- ▶ $\frac{72766}{2048} = 35 \text{ pages} + 1086 \text{ bytes}$
- ▶ Internal fragmentation: $2048 - 1086 = 962 \text{ bytes}$

Paging Example - Internal Fragmentation

- ▶ Page size = 2048 bytes
- ▶ Process size = 72766 bytes
- ▶ $\frac{72766}{2048} = 35 \text{ pages} + 1086 \text{ bytes}$
- ▶ Internal fragmentation: $2048 - 1086 = 962 \text{ bytes}$
- ▶ Worst case fragmentation = 1 frame - 1 byte

Paging Example - Internal Fragmentation

- ▶ Page size = 2048 bytes
- ▶ Process size = 72766 bytes
- ▶ $\frac{72766}{2048} = 35 \text{ pages} + 1086 \text{ bytes}$
- ▶ Internal fragmentation: $2048 - 1086 = 962 \text{ bytes}$
- ▶ Worst case fragmentation = 1 frame - 1 byte
- ▶ On average fragmentation = $\frac{1}{2}$ frame size



?

Small Page Size vs. Big Page Size

- ▶ On average fragmentation = $\frac{1}{2}$ page size, hence, small page sizes are desirable.

Small Page Size vs. Big Page Size

- ▶ On average fragmentation = $\frac{1}{2}$ page size, hence, small page sizes are desirable.
- ▶ Small pages, more overhead is in the page-table, this overhead is reduced as the size of the pages increases.

Small Page Size vs. Big Page Size

- ▶ On average fragmentation = $\frac{1}{2}$ page size, hence, small page sizes are desirable.
- ▶ Small pages, more overhead is in the page-table, this overhead is reduced as the size of the pages increases.
- ▶ Disk I/O is more efficient when the amount data being transferred is larger (e.g., big pages).

Small Page Size vs. Big Page Size

- ▶ On average fragmentation = $\frac{1}{2}$ page size, hence, small page sizes are desirable.
- ▶ Small pages, more overhead is in the page-table, this overhead is reduced as the size of the pages increases.
- ▶ Disk I/O is more efficient when the amount data being transferred is larger (e.g., big pages).
- ▶ Pages typically are between 4 KB and 8 KB in size.

```
getconf PAGESIZE
```


Page Table Implementation

Page Table

- ▶ Page table is kept in main memory.

Page Table

- ▶ Page table is kept in main memory.
- ▶ Page-table base register (PTBR) points to the page table.
- ▶ Page-table length register (PTLR) indicates size of the page table.

Page Table

- ▶ Page table is kept in main memory.
- ▶ Page-table base register (PTBR) points to the page table.
- ▶ Page-table length register (PTLR) indicates size of the page table.
- ▶ In this scheme every data/instruction access requires two memory accesses.

Page Table

- ▶ Page table is kept in main memory.
- ▶ Page-table base register (PTBR) points to the page table.
- ▶ Page-table length register (PTLR) indicates size of the page table.
- ▶ In this scheme every data/instruction access requires two memory accesses.
 - One for the page table and one for the data/instruction.

Translation Look-aside Buffers (TLB)

- ▶ The two memory access problem can be solved by the use of a special **fast-lookup hardware cache** called **associative memory** or **translation look-aside buffers (TLBs)**.

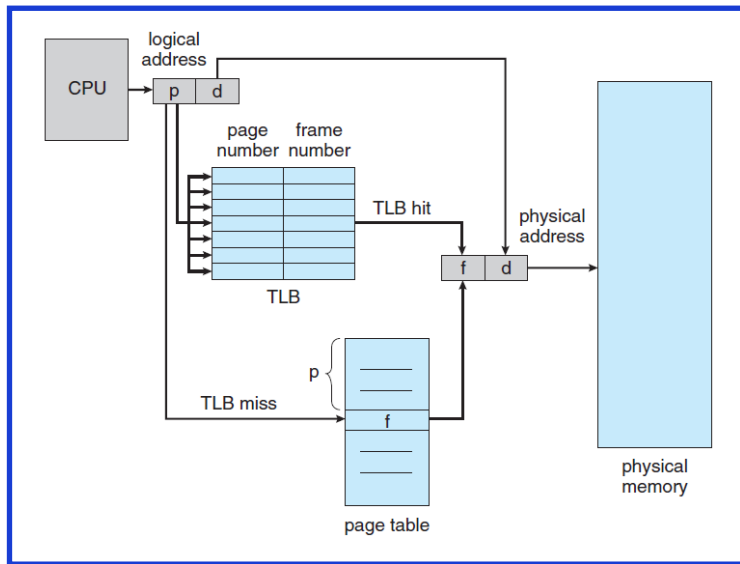
Associative Memory

- ▶ Associative memory: parallel search

Page #	Frame #

- ▶ Address translation (p, d)
 - If p is in associative register, get frame# out
 - Otherwise, get frame# from page table in memory

Paging Hardware With TLB



- ▶ **Hit ratio:** percentage of times that a page number is found in the TLB.

Effective Access Time

- ▶ **Hit ratio:** percentage of times that a page number is found in the TLB.
- ▶ **Effective Access Time (EAT)**
 - α : memory access latency
 - h : hit ratio
 - $EAT = h \times \alpha + (1 - h) \times 2\alpha$

Effective Access Time

- ▶ **Hit ratio:** percentage of times that a page number is found in the TLB.
- ▶ **Effective Access Time (EAT)**
 - α : memory access latency
 - h : hit ratio
 - $EAT = h \times \alpha + (1 - h) \times 2\alpha$
- ▶ $h = 80\%, \alpha = 100ns \Rightarrow EAT = 0.80 \times 100 + 0.20 \times 200 = 120ns$

Effective Access Time

- ▶ Hit ratio: percentage of times that a page number is found in the TLB.
- ▶ Effective Access Time (EAT)
 - α : memory access latency
 - h : hit ratio
 - $EAT = h \times \alpha + (1 - h) \times 2\alpha$
- ▶ $h = 80\%, \alpha = 100ns \Rightarrow EAT = 0.80 \times 100 + 0.20 \times 200 = 120ns$
- ▶ $h = 99\%, \alpha = 100ns \Rightarrow EAT = 0.99 \times 100 + 0.01 \times 200 = 101ns$

More About TLB

- ▶ Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry
 - **Uniquely identifies each process** to provide address-space protection for that process.
 - Otherwise, need to **flush** at every **context switch**.

More About TLB

- ▶ Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry
 - **Uniquely identifies each process** to provide address-space protection for that process.
 - Otherwise, need to **flush** at every **context switch**.
- ▶ TLBs typically **small** (64 to 1,024 entries)

More About TLB

- ▶ Some TLBs store address-space identifiers (ASIDs) in each TLB entry
 - Uniquely identifies each process to provide address-space protection for that process.
 - Otherwise, need to flush at every context switch.
- ▶ TLBs typically small (64 to 1,024 entries)
- ▶ On a TLB miss, value is loaded into the TLB for faster access next time.
 - Replacement policies must be considered.

?

Memory Protection

- ▶ Memory protection implemented by associating **protection bit** with each frame to indicate if **read-only** or **read-write** access is allowed.


Memory Protection

- ▶ Memory protection implemented by associating **protection bit** with each frame to indicate if **read-only** or **read-write** access is allowed.
- ▶ **Valid-invalid bit** attached to **each entry** in the page table:
 - **Valid** indicates that the associated page is **in the process logical address space**, and is thus a legal page.
 - **Invalid** indicates that the page is **not in the process logical address space**.

Memory Protection

- ▶ Memory protection implemented by associating **protection bit** with each frame to indicate if **read-only** or **read-write** access is allowed.
- ▶ **Valid-invalid bit** attached to **each entry** in the page table:
 - **Valid** indicates that the associated page is **in the process logical address space**, and is thus a legal page.
 - **Invalid** indicates that the page is **not in the process logical address space**.
 - Or use **page-table length register (PTLR)**.

Memory Protection

- ▶ Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed.
- ▶ Valid-invalid bit attached to each entry in the page table:
 -  • Valid indicates that the associated page is in the process logical address space, and is thus a legal page.
 - Invalid indicates that the page is not in the process logical address space.
 - Or use page-table length register (PTLR).
- ▶ Any violations result in a trap to the kernel.

Valid/Invalid Bit In A Page Table

00000	page 0
	page 1
	page 2
	page 3
	page 4
10,468	page 5
12,287	

frame number		valid-invalid bit	
0	2	v	
1	3	v	
2	4	v	
3	7	v	
4	8	v	
5	9	v	
6	0	i	
7	0	i	

page table

0	
1	
2	page 0
3	page 1
4	page 2
5	
6	
7	page 3
8	page 4
9	page 5
	⋮
	page <i>n</i>

Shared Pages

► Shared code

- One copy of **read-only (reentrant) code** shared among processes (e.g., text editors).
- Similar to multiple threads sharing the same process space.

Shared Pages

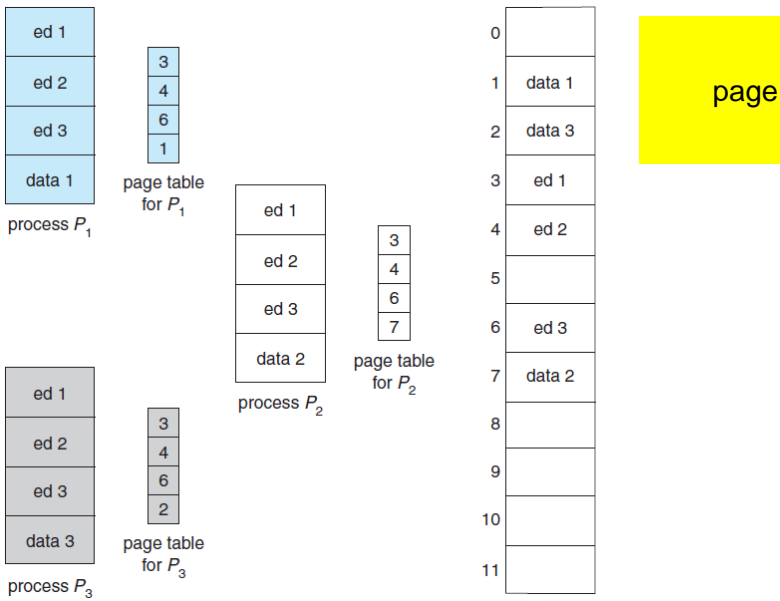
► Shared code

- One copy of read-only (reentrant) code shared among processes (e.g., text editors).
- Similar to multiple threads sharing the same process space.

► Private code and data

- Each process keeps a separate copy of the code and data.
- The pages for the private code and data can appear anywhere in the logical address space.

Shared Pages Example



Structure of the Page Table

Structure of the Page Table (1/2)

- ▶ Memory structures for paging can get huge using straight-forward methods.

Structure of the Page Table (1/2)

- ▶ Memory structures for paging can get huge using straight-forward methods.
- ▶ Consider a 32-bit logical address space as on modern computers:

Structure of the Page Table (1/2)

- ▶ Memory structures for paging can get huge using straight-forward methods.
- ▶ Consider a 32-bit logical address space as on modern computers:
 - Page size of $4KB = 2^{12}$.

Structure of the Page Table (1/2)

- ▶ Memory structures for paging can get huge using straight-forward methods.
- ▶ Consider a 32-bit logical address space as on modern computers:
 - Page size of $4KB = 2^{12}$.
 - Page table would have 1 million entries $\left(\frac{2^{32}}{2^{12}}\right)$.

Structure of the Page Table (1/2)

- ▶ Memory structures for paging can get huge using straight-forward methods.
- ▶ Consider a 32-bit logical address space as on modern computers:
 - Page size of $4KB = 2^{12}$.
 - Page table would have 1 million entries ($\frac{2^{32}}{2^{12}}$).
 - If each entry is 4B: 4MB of physical address space memory for page table alone.

Structure of the Page Table (1/2)

- ▶ Memory structures for paging can get huge using straight-forward methods.
- ▶ Consider a 32-bit logical address space as on modern computers:
 - Page size of $4KB = 2^{12}$.
 - Page table would have 1 million entries $\left(\frac{2^{32}}{2^{12}}\right)$.
 - If each entry is 4B: 4MB of physical address space memory for page table alone.
 - That amount of memory used to cost a lot.

Structure of the Page Table (1/2)

- ▶ Memory structures for paging can get huge using straight-forward methods.
- ▶ Consider a 32-bit logical address space as on modern computers:
 - Page size of $4KB = 2^{12}$.
 - Page table would have 1 million entries $\left(\frac{2^{32}}{2^{12}}\right)$.
 - If each entry is 4B: 4MB of physical address space memory for page table alone.
 - That amount of memory used to cost a lot.
 - Don't want to allocate that contiguously in main memory.

Structure of the Page Table (2/2)


- ▶ Hierarchical Paging
- ▶ Hashed Page Tables
- ▶ Inverted Page Tables

Hierarchical Paging

Hierarchical Page Tables

- ▶ Break up the logical address space into multiple page tables.
- ▶ A simple technique is a two-level page table.
- ▶ We then page the page table.

Two-Level Page-Table Scheme



Two-Level Paging Example

- ▶ A logical address, on 32-bit machine with 1K page size, is divided:

Two-Level Paging Example

- ▶ A logical address, on 32-bit machine with 1K page size, is divided:
 - A page number consisting of 22 bits.

Two-Level Paging Example

- ▶ A logical address, on 32-bit machine with 1K page size, is divided:
 - A page number consisting of 22 bits.
 - A page offset consisting of 10 bits.


Two-Level Paging Example

- ▶ A logical address, on 32-bit machine with 1K page size, is divided:
 - A page number consisting of 22 bits.
 - A page offset consisting of 10 bits.
- ▶ Since the page table is paged, the page number is divided into:
 - A 12-bit page number.
 - A 10-bit page offset.

Two-Level Paging Example

- ▶ A logical address, on 32-bit machine with 1K page size, is divided:
 - A page number consisting of 22 bits.
 - A page offset consisting of 10 bits.
- ▶ Since the page table is paged, the page number is divided into:
 - A 12-bit page number.
 - A 10-bit page offset.

- ▶ Thus, a logical address is:




- ▶ where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table.

Two-Level Paging Example

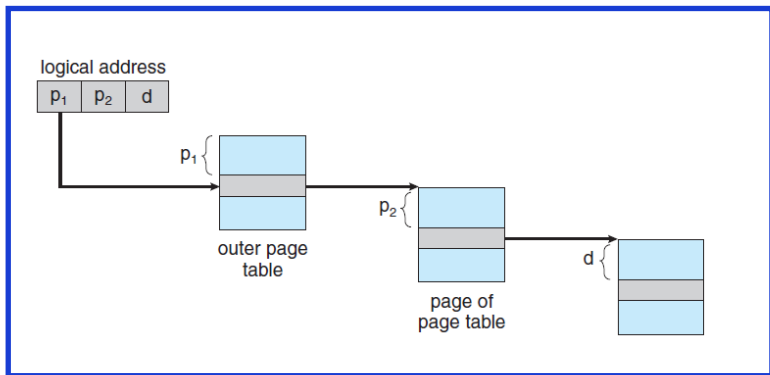
- ▶ A logical address, on 32-bit machine with 1K page size, is divided:
 - A page number consisting of 22 bits.
 - A page offset consisting of 10 bits.
- ▶ Since the page table is paged, the page number is divided into:
 - A 12-bit page number.
 - A 10-bit page offset.

- ▶ Thus, a logical address is:



- ▶ where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table.
- ▶ Known as forward-mapped page table.

Address-Translation Scheme



64-bit Logical Address Space

- ▶ Even two-level paging scheme **not sufficient**.
- ▶ If page size is **4KB (2^{12})**

64-bit Logical Address Space

- ▶ Even two-level paging scheme **not sufficient**.
- ▶ If page size is **4KB (2^{12})**
 - Then page table has **2^{52}** entries

64-bit Logical Address Space

- ▶ Even two-level paging scheme **not sufficient**.
- ▶ If page size is **4KB (2^{12})**
 - Then page table has **2^{52}** entries
 - If two level scheme, **inner page tables** could be **2^{10} , 4B** entries

64-bit Logical Address Space

- ▶ Even two-level paging scheme **not sufficient**.
- ▶ If page size is **4KB (2^{12})**
 - Then page table has **2^{52}** entries
 - If two level scheme, **inner page tables** could be **2^{10} , 4B** entries
 - **Outer page table** has **2^{42}** entries or **$2^{44}B$**

64-bit Logical Address Space

- ▶ Even two-level paging scheme **not sufficient**.
- ▶ If page size is **4KB (2^{12})**
 - Then page table has 2^{52} entries
 - If two level scheme, **inner page tables** could be 2^{10} , **4B** entries
 - **Outer page table** has 2^{42} entries or $2^{44}B$
 - Address would look like:

64-bit Logical Address Space

- ▶ Even two-level paging scheme not sufficient.
- ▶ If page size is 4KB (2^{12})
 - Then page table has 2^{52} entries
 - If two level scheme, inner page tables could be 2^{10} , 4B entries
 - Outer page table has 2^{42} entries or $2^{44}B$
 - Address would look like:

outer page	inner page	offset
p_1	p_2	d
42	10	12

Three-level Paging Scheme

- ▶ One solution is to add a 2nd outer page table.

Three-level Paging Scheme

- ▶ One solution is to add a 2nd outer page table.
- ▶ But in the following example the 2nd outer page table is still 2^{34} bytes in size.

Three-level Paging Scheme

- ▶ One solution is to add a 2nd outer page table.
- ▶ But in the following example the 2nd outer page table is still 2^{34} bytes in size.
- ▶ And possibly 4 memory access to get to one physical memory location.

outer page	inner page	offset
p_1	p_2	d
42	10	12

2nd outer page	outer page	inner page	offset
p_1	p_2	p_3	d
32	10	10	12

Hashed Page Tables

Hashed Page Tables (1/2)

- ▶ Common in address spaces > 32 bits
- ▶ The logical page number is hashed into a page table.
- ▶ This page table contains a chain of elements hashing to the same location.

Hashed Page Tables (2/2)

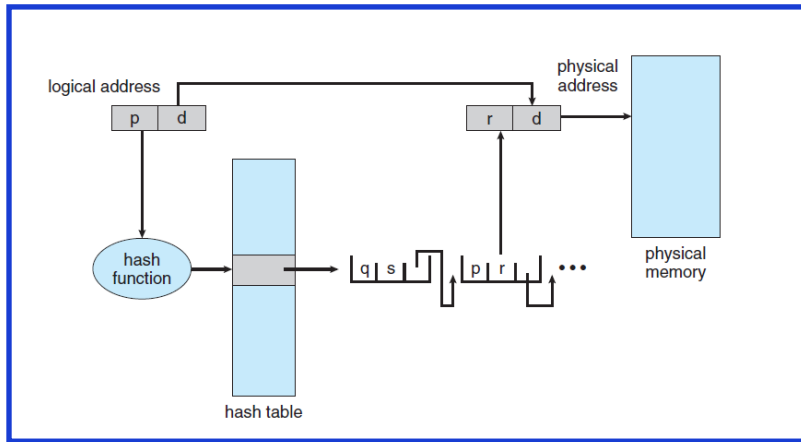
- ▶ Each element contains
 - ① The **logical page number**
 - ② The value of the **mapped page frame**
 - ③ A **pointer** to the next element

Hashed Page Tables (2/2)

- ▶ Each element contains
 - ① The logical page number
 - ② The value of the mapped page frame
 - ③ A pointer to the next element
- ▶ Logical page numbers are compared in this chain searching for a match.
 - If a match is found, the corresponding physical frame is extracted.



Hashed Page Table Architecture



Inverted Page Tables

Inverted Page Table (1/2)

- ▶ Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages.

Inverted Page Table (1/2)

- ▶ Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages.
- ▶ One entry for each real page of memory.

Inverted Page Table (1/2)

- ▶ Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages.
- ▶ One entry for each real page of memory.
- ▶ Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page.

Inverted Page Table (2/2)

- ▶ Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs.

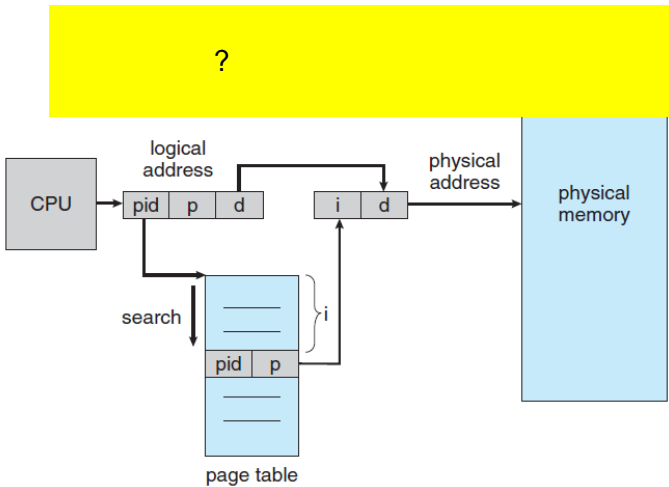
Inverted Page Table (2/2)

- ▶ Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs.
- ▶ Use hash table to limit the search to one, or at most a few, page-table entries.

Inverted Page Table (2/2)

- ▶ Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs.
- ▶ Use hash table to limit the search to one, or at most a few, page-table entries.
- ▶ But how to implement shared memory?
 - One mapping of a virtual address to the shared physical address

Inverted Page Table Architecture



Summary

- ▶ Paging vs. Segmentation

Summary

- ▶ Paging vs. Segmentation
- ▶ Physical memory: frames, Logical memory: pages

Summary

- ▶ Paging vs. Segmentation
- ▶ Physical memory: frames, Logical memory: pages
- ▶ Page table: translates logical to physical addresses

- ▶ Paging vs. Segmentation
- ▶ Physical memory: frames, Logical memory: pages
- ▶ Page table: translates logical to physical addresses
- ▶ Translation Look-aside Buffer (TLB)

- ▶ Paging vs. Segmentation
- ▶ Physical memory: frames, Logical memory: pages
- ▶ Page table: translates logical to physical addresses
- ▶ Translation Look-aside Buffer (TLB)
- ▶ Memory protection: valid-invalid bit

Summary

- ▶ Paging vs. Segmentation
- ▶ Physical memory: frames, Logical memory: pages
- ▶ Page table: translates logical to physical addresses
- ▶ Translation Look-aside Buffer (TLB)
- ▶ Memory protection: valid-invalid bit
- ▶ Page table structure: hierarchical paging, hashed page tables, inverted page tables

Questions?

Acknowledgements

Some slides were derived from Avi Silberschatz slides.