

Virtual Memory (Part I)

Amir H. Payberah
amir@sics.se

Amirkabir University of Technology
(Tehran Polytechnic)



Motivation and Background

Motivation

- ▶ A **program** needs to be in **memory** to execute.
- ▶ But, **entire program rarely used**.
 - Error code, unusual routines, large data structures
- ▶ Entire program code **not needed at same time**.

- ▶ Consider ability to execute **partially-loaded program**.

Motivation

- ▶ Consider ability to execute **partially-loaded program**.
- ▶ Program **no longer constrained** by limits of **physical memory**.

Motivation

- ▶ Consider ability to execute **partially-loaded program**.
- ▶ Program **no longer constrained** by limits of **physical memory**.
- ▶ Each program takes **less memory** while running: **more programs run at the same time**.
 - **Increased CPU utilization** and **throughput** with no increase in response time or turnaround time.

Motivation

- ▶ Consider ability to execute partially-loaded program.
- ▶ Program no longer constrained by limits of physical memory.
- ▶ Each program takes less memory while running: more programs run at the same time.
 - Increased CPU utilization and throughput with no increase in response time or turnaround time.
- ▶ Less I/O needed to load or swap programs into memory: each user program runs faster.

Virtual Memory

Virtual Memory (1/2)

- ▶ Separation of user logical memory from physical memory.

Virtual Memory (1/2)

- ▶ Separation of user logical memory from physical memory.
- ▶ Only part of the program needs to be in memory for execution.


Virtual Memory (1/2)

- ▶ Separation of user logical memory from physical memory.
- ▶ Only part of the program needs to be in memory for execution.
- ▶ Logical address space can therefore be much larger than physical address space.

Virtual Memory (1/2)

- ▶ Separation of user logical memory from physical memory.
- ▶ Only part of the program needs to be in memory for execution.
- ▶ Logical address space can therefore be much larger than physical address space.
- ▶ More programs running concurrently.

Virtual Memory (2/2)



Virtual Address Space (1/3)

- ▶ **Virtual address space**: **logical view** of how process is stored in memory.

Virtual Address Space (1/3)

- ▶ **Virtual address space**: **logical view** of how process is stored in memory.
- ▶ Usually **starts at address 0**, contiguous addresses **until end of space**.


Virtual Address Space (1/3)

- ▶ **Virtual address space**: **logical view** of how process is stored in memory.
- ▶ Usually **starts at address 0**, contiguous addresses **until end of space**.
- ▶ Meanwhile, **physical memory** organized in page **frames**.

Virtual Address Space (1/3)

- ▶ Virtual address space: logical view of how process is stored in memory.
- ▶ Usually starts at address 0, contiguous addresses until end of space.
- ▶ Meanwhile, physical memory organized in page frames.
- ▶ MMU must map logical to physical.

Virtual Address Space (2/3)



Virtual Address Space (3/3)

- ▶ The **heap** grows upward in memory, and the **stack** grows downward.

Virtual Address Space (3/3)

- ▶ The **heap** grows upward in memory, and the **stack** grows downward.
- ▶ The **hole** between the heap and the stack is part of the **virtual address space**, but will require actual physical pages only if the heap or stack **grows**.

Virtual Address Space (3/3)




- ▶ The **heap** grows upward in memory, and the **stack** grows downward.
- ▶ The **hole** between the heap and the stack is part of the **virtual address space**, but will require actual physical pages only if the heap or stack **grows**.
- ▶ Virtual address spaces that **include holes** are known as **sparse address spaces**.

Heap vs. Stack

► Stack

- Stores **local variables** created by functions.
- New variables are **pushed** onto the stack.
- When a function exits, all of the its pushed variables are **freed**.




Heap vs. Stack

► Stack

- Stores local variables created by functions.
- New variables are **pushed** onto the stack.
- When a function exits, all of the its pushed variables are freed.


► Heap

- Used for **dynamic allocation**.
- Use `malloc()` or `calloc()` to allocate memory on the heap.
- Use `free()` to deallocate the memory.



Virtual Memory Benefits


- ▶ Separating logical memory from physical memory.
- ▶ Page sharing.



Demand Paging


Demand Paging (1/2)

- Could bring **entire process** into memory at load time, or bring **a page** into memory only **when it is needed**.




Demand Paging (1/2)

- ▶ Could bring entire process into memory at load time, or bring a page into memory only when it is needed.
- ▶ A demand-paging system is similar to a paging system with swapping.




Demand Paging (2/2)

- ▶ Rather than swapping the **entire process** into memory, we use a **lazy swapper**.
- ▶ A **lazy swapper** never swaps a page into memory unless that page will be **needed**.



Demand Paging (2/2)

- ▶ Rather than swapping the entire process into memory, we use a lazy swapper.
- ▶ A lazy swapper never swaps a page into memory unless that page will be needed.
- ▶ A swapper that deals with pages is a pager.



Basic Concepts




- ▶ The **pager** guesses **which pages** will be used before swapping out again.

Basic Concepts

- ▶ The **pager** guesses **which pages** will be used before swapping out again.
- ▶ The pager brings only **the needed pages** into memory.

Basic Concepts

- ▶ The pager guesses which pages will be used before swapping out again.
- ▶ The pager brings only the needed pages into memory.
- ▶ Uses valid-invalid bit to distinguish between the pages that are in memory and the pages that are on the disk?
v: memory resident
i: not in memory

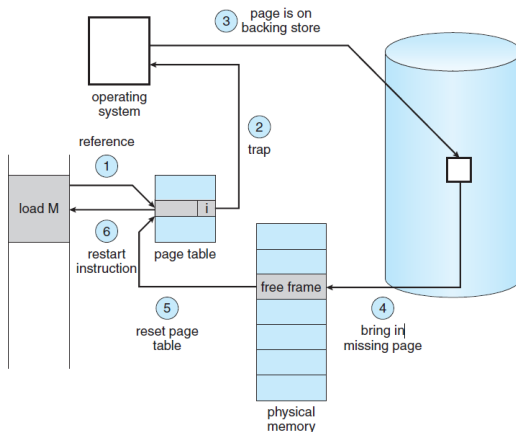


Page Fault

- ▶ Access to a page marked **invalid** causes a **page fault**.
- ▶ Causing a trap to the OS: brings the **desired page into memory**.

Handling Page Fault (1/6)

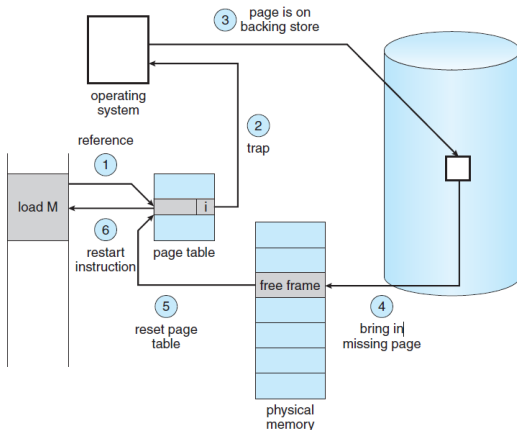
- Check an internal table for the process to determine whether the reference was a valid or an invalid memory access.



Handling Page Fault (2/6)

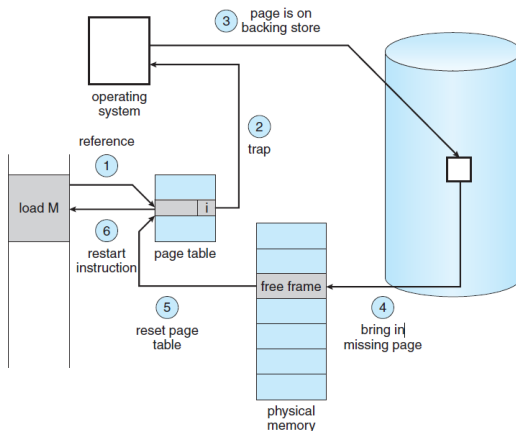


- ▶ If the reference was invalid, we terminate the process.
- ▶ If it was valid but we have not yet brought in that page, we now page it in.



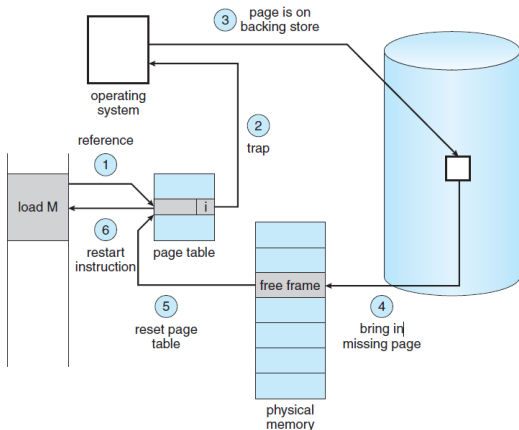
Handling Page Fault (3/6)

- We find a free frame.




Handling Page Fault (4/6)

- ▶ We schedule a disk operation to read the desired page into the newly allocated frame.



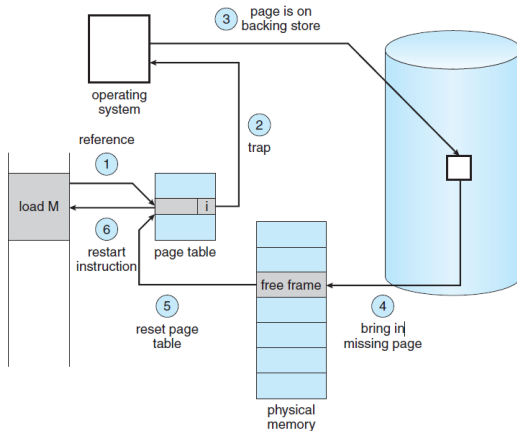
Handling Page Fault (5/6)

- ▶ When the disk read is complete, we **modify** the **internal table** kept with the process and the page table to indicate that the page is now in memory.



Handling Page Fault (6/6)

- ▶ We restart the instruction that was interrupted by the trap.



Aspects of Demand Paging (1/2)

- ▶ Extreme case: start process with no pages in memory.

Aspects of Demand Paging (1/2)

- ▶ Extreme case: start process with no pages in memory.
- ▶ OS sets instruction pointer to first instruction of process, non-memory-resident: page fault

Aspects of Demand Paging (1/2)

- ▶ Extreme case: start process with no pages in memory.
- ▶ OS sets instruction pointer to first instruction of process, non-memory-resident: page fault
- ▶ And for every other process pages on first access.

Aspects of Demand Paging (1/2)

- ▶ Extreme case: start process with no pages in memory.
- ▶ OS sets instruction pointer to first instruction of process, non-memory-resident: page fault
- ▶ And for every other process pages on first access.
- ▶ Pure demand paging

Aspects of Demand Paging (2/2)

- ▶ A program could access **several new pages** of memory with each instruction execution: **multiple page faults** per instruction.

Aspects of Demand Paging (2/2)

- ▶ A program could access **several new pages** of memory with each instruction execution: **multiple page faults** per instruction.
- ▶ **Unacceptable system performance**

Aspects of Demand Paging (2/2)

- ▶ A program could access several new pages of memory with each instruction execution: multiple page faults per instruction.
- ▶ Unacceptable system performance
- ▶ **?** Locality of reference results in reasonable performance from demand paging.

Demand Paging Hardware



- ▶ The hardware to support demand paging is the same as the hardware for paging and swapping:
 - Page table with valid-invalid bit
 - Secondary memory with swap space
- ▶ The ability to restart any instruction after a page fault.

Stages in Demand Paging (1/3)

- ▶ 1. **Trap** to the OS.

Stages in Demand Paging (1/3)

- ▶ 1. **Trap** to the OS.
- ▶ 2. Save the **user registers** and **process state**.

Stages in Demand Paging (1/3)

- ▶ 1. **Trap** to the OS.
- ▶ 2. Save the **user registers** and **process state**.
- ▶ 3. Determine that the **interrupt** was a **page fault**.

Stages in Demand Paging (1/3)

- ▶ 1. Trap to the OS.
- ▶ 2. Save the user registers and process state.
- ▶ 3. Determine that the interrupt was a page fault.
- ▶ 4. Check that the page reference was legal and determine the location of the page on the disk.

Stages in Demand Paging (2/3)

- ▶ 5. Issue a **read from the disk** to a **free frame**:
 - **Wait** in a **queue** for this device until the read request is serviced.
 - **Wait** for the **device seek and/or latency time**.
 - Begin the **transfer** of the page to a free frame.

Stages in Demand Paging (2/3)

- ▶ 5. Issue a **read from the disk** to a **free frame**:
 - **Wait** in a **queue** for this device until the read request is serviced.
 - **Wait** for the **device seek and/or latency time**.
 - Begin the **transfer** of the page to a free frame.

- ▶ 6. While waiting, **allocate the CPU to some other user**.

Stages in Demand Paging (2/3)

- ▶ 5. Issue a read from the disk to a free frame:
 - Wait in a queue for this device until the read request is serviced.
 - Wait for the device seek and/or latency time.
 - Begin the transfer of the page to a free frame.
- ▶ 6. While waiting, allocate the CPU to some other user.
- ▶ 7. Receive an interrupt from the disk I/O subsystem (I/O completed).

Stages in Demand Paging (3/3)

- ▶ 8. Save the **registers and process state** for the other user.

Stages in Demand Paging (3/3)

- ▶ 8. Save the **registers and process state** for the other user.
- ▶ 9. Determine that the **interrupt** was from the **disk**.

Stages in Demand Paging (3/3)

- ▶ 8. Save the **registers and process state** for the other user.
- ▶ 9. Determine that the **interrupt** was from the **disk**.
- ▶ 10. Correct the **page table** and other tables to show page is now in **memory**.

Stages in Demand Paging (3/3)

- ▶ 8. Save the **registers and process state** for the other user.
- ▶ 9. Determine that the **interrupt** was from the **disk**.
- ▶ 10. Correct the **page table** and other tables to show page is now in **memory**.
- ▶ 11. **Wait** for the **CPU to be allocated** to this process again.

Stages in Demand Paging (3/3)

- ▶ 8. Save the registers and process state for the other user.
- ▶ 9. Determine that the interrupt was from the disk.
- ▶ 10. Correct the page table and other tables to show page is now in memory.
- ▶ 11. Wait for the CPU to be allocated to this process again.
- ▶ 12. Restore the user registers, process state, and new page table, and then resume the interrupted instruction.

Performance of Demand Paging

- ▶ Probability of a page fault: $0 \leq p \leq 1$

Performance of Demand Paging

- ▶ Probability of a page fault: $0 \leq p \leq 1$
- ▶ If $p = 0$: no page faults
- ▶ If $p = 1$: every reference is a fault

Performance of Demand Paging

- ▶ Probability of a page fault: $0 \leq p \leq 1$
- ▶ If $p = 0$: no page faults
- ▶ If $p = 1$: every reference is a fault
- ▶ Effective Access Time (EAT)
- ▶ $EAT = (1 - p) \times \text{memory_access} + p \times \text{page_fault_time}$

Demand Paging Example

- ▶ Memory access time = 200 nanoseconds
- ▶ Average page-fault service time = 8 milliseconds

Demand Paging Example

- ▶ Memory access time = 200 nanoseconds
- ▶ Average page-fault service time = 8 milliseconds
- ▶ $EAT = (1 - p) \times 200 + p \times 8,000,000 = 200 + p \times 7,999,800$

Demand Paging Example

- ▶ Memory access time = 200 nanoseconds
- ▶ Average page-fault service time = 8 milliseconds
- ▶ $EAT = (1 - p) \times 200 + p \times 8,000,000 = 200 + p \times 7,999,800$
- ▶ If $p = 0.001$, then $EAT = 8.2 \text{ microseconds}$: this is a slowdown by a factor of 40.

Demand Paging Example

- ▶ Memory access time = 200 nanoseconds
- ▶ Average page-fault service time = 8 milliseconds
- ▶ $EAT = (1 - p) \times 200 + p \times 8,000,000 = 200 + p \times 7,999,800$
- ▶ If $p = 0.001$, then $EAT = 8.2 \text{ microseconds}$: this is a slowdown by a factor of 40.
- ▶ If want performance degradation < 10 percent:
 - $220 > 200 + p \times 7,999,800$
 - $20 > p \times 7,999,800$
 - $p < .0000025$: less than one page fault in every 400,000 memory accesses

Demand Paging Optimization

- ▶ Swap space I/O faster than file system I/O even if on the same device.
 - Swap allocated in larger chunks, less management needed than file system.

Demand Paging Optimization

- ▶ Swap space I/O faster than file system I/O even if on the same device.
 - Swap allocated in larger chunks, less management needed than file system.
- ▶ Copy entire process image to swap space at process load time.
 - Then page in and out of swap space.

Demand Paging Optimization

- ▶ Swap space I/O faster than file system I/O even if on the same device.
 - Swap allocated in larger chunks. less management needed than file system.
- ▶ Copy entire process image to swap space at process load time.
 - Then page in and out of swap space.
- ▶ Demand page in from program binary on disk, but discard rather than paging out when freeing frame.
 - Pages not associated with a file, i.e., heap and stack (anonymous memory) still need to write to swap space.


?

Copy-on-Write

Copy-on-Write


- ▶ Copy-on-Write allows both **parent** and **child** processes to initially **share the same pages in memory**.
- ▶ If either process **modifies a shared page**, only then is the page copied.

Copy-on-Write Example




Before modification

Copy-on-Write Example



Before modification



After modification

More About Copy-on-Write

- ▶ **Copy-on-Write** allows **more efficient process creation** as only modified pages are copied.

More About Copy-on-Write

- ▶ **Copy-on-Write** allows **more efficient process creation** as only modified pages are copied.
- ▶ Free pages are allocated from a pool of **zero-fill-on-demand** pages.
 - **Pool** should always have **free frames** for fast demand page execution.
 - **Zero-fill-on-demand** pages have been zeroed-out before being allocated, thus **erasing the previous contents**.

More About Copy-on-Write

- ▶ Copy-on-Write allows more efficient process creation as only modified pages are copied.
- ▶ Free pages are allocated from a pool of **zero-fill-on-demand** pages.
 - **Pool** should always have **free frames** for fast demand page execution.
 - **Zero-fill-on-demand** pages have been zeroed-out before being allocated, thus **erasing the previous contents**.
- ▶ **vfork()** variation on **fork()** does not use Copy-on-Write: with **vfork()**, the parent process is suspended, and the child process uses the address space of the parent.

?

Page Replacement

What Happens if There is no Free Frame?

- ▶ Assume, we had **forty frames** in **physical memory**.
- ▶ And, we run **six processes**, each of which is **ten pages** in size, but actually uses only **five pages**.

What Happens if There is no Free Frame?

- ▶ Assume, we had **forty frames** in **physical memory**.
- ▶ And, we run **six processes**, each of which is **ten pages** in size, but actually uses only **five pages**.
- ▶ It is possible that each of these processes may **suddenly try to use all ten of its pages**: resulting in a need for **sixty frames** when only **forty** are available.

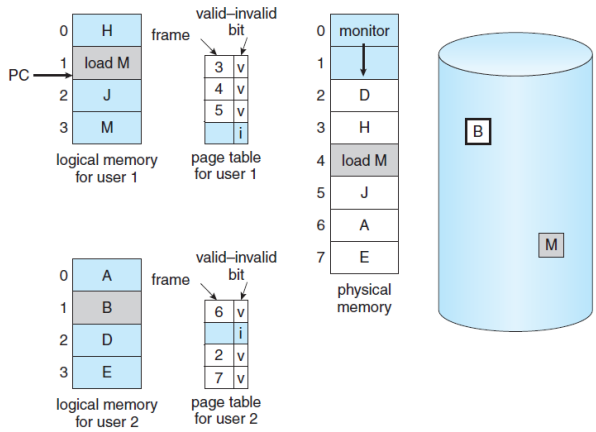
What Happens if There is no Free Frame?

- ▶ Assume, we had **forty frames** in **physical memory**.
- ▶ And, we run **six processes**, each of which is **ten pages** in size, but actually uses only **five pages**.
- ▶ It is possible that each of these processes may **suddenly try to use all ten of its pages**: resulting in a need for **sixty frames** when only **forty** are available.
- ▶ **Increasing the degree of multiprogramming: over-allocating memory**

Over-Allocation of Memory

- ▶ While a user process is **executing**, a page fault occurs.
- ▶ The OS determines where the desired page is residing on the **disk**.
- ▶ But, it finds that there are no free frames on the free-frame list.
- ▶ Need for page replacement

Need For Page Replacement



Basic Page Replacement

- 1 Find the location of the desired page on disk.

Basic Page Replacement

- ① Find the location of the desired page on disk.
- ② Find a free frame.
 - If there is a free frame, use it.
 - If there is no free frame, use a page replacement algorithm to select a victim frame: write victim frame to disk if dirty.


Basic Page Replacement

- ① Find the location of the desired page on disk.
- ② Find a free frame.
 - If there is a free frame, use it.
 - If there is no free frame, use a page replacement algorithm to select a victim frame: write victim frame to disk if dirty.
- ③ Bring the desired page into the (newly) free frame; update the page and frame tables

Basic Page Replacement


- ① Find the location of the desired page on disk.
- ② Find a free frame.
 - If there is a free frame, use it.
 - If there is no free frame, use a page replacement algorithm to select a victim frame: write victim frame to disk if dirty.
- ③ Bring the desired page into the (newly) free frame; update the page and frame tables
- ④ Continue the process by restarting the instruction that caused the trap.

Page Replacement



- Use modify (dirty) bit to reduce overhead of page transfers - only modified pages are written to disk.

Page Faults vs. The Number of Frames



Evaluate Page Replacement Algorithms

- ▶ Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string.
- ▶ String is just page numbers, not full addresses.
- ▶ For example, a reference string could be
7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1


Page Replacement Algorithms

- ▶ First-In-First-Out (FIFO) page replacement
- ▶ Optimal page replacement
- ▶ Least Recently Used (LRU) page replacement
- ▶ LRU-Approximation page replacement
- ▶ Counting-Based page replacement

FIFO Page Replacement

FIFO Page Replacement

- ▶ Reference string: 7,0,1,2,0,3,0,4,2,3,0,3,2,1,2,0,1,7,0,1
- ▶ 3 frames (3 pages can be in memory at a time per process)



FIFO Page Replacement

- ▶ Reference string: 7,0,1,2,0,3,0,4,2,3,0,3,2,1,2,0,1,7,0,1
- ▶ 3 frames (3 pages can be in memory at a time per process)

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1


7	7	7	2																
	0	0	0		2	2	4	4	4	0			0	0			7	7	7
					3	3	3	2	2	2			1	1			1	0	0
		1	1		1	0	0	0	3	3			3	2			2	2	1

page frames

- ▶ 15 page faults

FIFO Belady's Anomaly

- ▶ Adding more frames can cause more page faults: Belady's Anomaly




Optimal Page Replacement

Optimal Page Replacement

- ▶ Replace page that will not be used for **longest period of time**: 9 page fault is **optimal** for the example.
- ▶ How do you know this? Can't read the future

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1




page frames

Optimal Page Replacement

- ▶ Replace page that will not be used for **longest period of time**: 9 page fault is **optimal** for the example.
- ▶ How do you know this? Can't read the future

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

- ▶ Used for measuring **how well your algorithm performs**.

LRU Page Replacement

LRU Page Replacement

- ▶ Use **past knowledge** rather than the **future**.
- ▶ Replace page that has **not been used in the most amount of time**

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0		1		1		1
	0	0	0		0		0	0	3	3		3		0		0
		1	1		3		3	2	2	2		2		2		7

page frames

LRU Page Replacement

- ▶ Use **past knowledge** rather than the **future**.
- ▶ Replace page that has **not been used in the most amount of time**

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0		1		1		1
	0	0	0		0		0	0	3	3		3		0		0
		1	1		3		3	2	2	2		2		2		7

page frames

- ▶ **12 faults: better than FIFO but worse than OPT**

LRU Page Replacement

- ▶ Use **past knowledge** rather than the **future**.
- ▶ Replace page that has **not been used in the most amount of time**

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames

- ▶ 12 faults: better than FIFO but worse than OPT
- ▶ Generally good algorithm and frequently used

LRU Implementation (1/2)

- ▶ Counter implementation
- ▶ Every page entry has a **counter**; every time page is referenced through this entry, copy the **clock into the counter**.
- ▶ When a page needs to be changed, look at the counters to find **smallest value**.
- ▶ Search through table needed.

LRU Implementation (2/2)

- ▶ Stack implementation
- ▶ Keep a stack of page numbers in a double link form.

LRU Implementation (2/2)

- ▶ Stack implementation
- ▶ Keep a stack of page numbers in a double link form.
- ▶ Page referenced:
 - Move it to the top
 - Requires 6 pointers to be changed

LRU Implementation (2/2)


- ▶ Stack implementation
- ▶ Keep a stack of page numbers in a double link form.
- ▶ Page referenced:
 - Move it to the top
 - Requires 6 pointers to be changed
- ▶ No search for replacement.

LRU Implementation (2/2)

- ▶ Stack implementation
- ▶ Keep a stack of page numbers in a double link form.
- ▶ Page referenced:
 - Move it to the top
 - Requires 6 pointers to be changed
- ▶ No search for replacement.
- ▶ LRU and OPT are cases of stack algorithms with no Belady's Anomaly.

Stack Implementation

- Use of a stack to record **most recent page references**.



LRU-Approximation Page Replacement

LRU-Approximation Page Replacement

- ▶ LRU needs special hardware and still slow
- ▶ Improvements: LRU-Approximation
 - Reference bit
 - Second-chance algorithm
 - Enhanced second-chance algorithm

- ▶ With each page associate a bit, initially = 0

Reference Bit

- ▶ With each page associate a bit, initially = 0
- ▶ When page is referenced, bit set to 1

Reference Bit

- ▶ With each page associate a bit, initially = 0
- ▶ When page is referenced, bit set to 1
- ▶ Replace any with reference bit = 0 (if one exists)

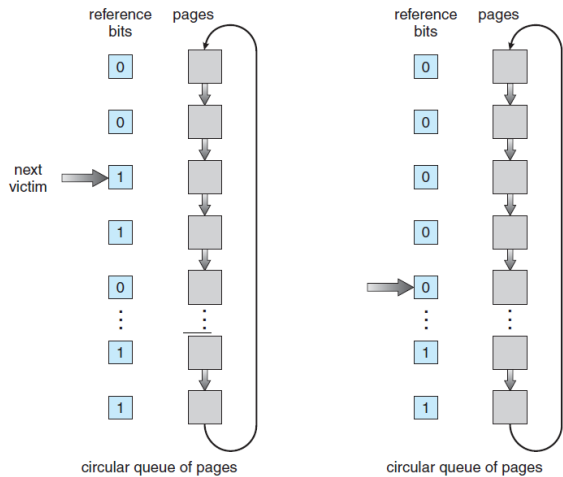
Reference Bit

- ▶ With each page associate a bit, initially = 0
- ▶ When page is referenced, bit set to 1
- ▶ Replace any with reference bit = 0 (if one exists)
- ▶ We do not know the order

Second-Chance Algorithm (1/2)

- ▶ It is also called **clock algorithm**.
- ▶ Generally **FIFO**, plus hardware-provided **reference bit**
- ▶ **If page to be replaced has**
 - Reference bit = 0 → **replace it**
 - Reference bit = 1 then, set reference bit 0, **leave page in memory**, and replace **next page**, subject to **same rules**.

Second-Chance Algorithm (2/2)



Enhanced Second-Chance Algorithm

- ▶ Improve algorithm by using **reference bit** and **modify bit**
(reference, modify)

Enhanced Second-Chance Algorithm

- ▶ Improve algorithm by using **reference bit** and **modify bit**
(**reference**, **modify**)
- ▶ Take ordered pair:
 - ① (0, 0) neither recently used nor modified: **best page** to replace
 - ② (0, 1) not recently used but modified: not quite as good, **must write out before replacement**
 - ③ (1, 0) recently used but clean: **probably will be used again soon**
 - ④ (1, 1) recently used and modified: **probably will be used again soon** and need to write out before replacement

Enhanced Second-Chance Algorithm

- ▶ Improve algorithm by using **reference bit** and **modify bit** (**reference, modify**)
- ▶ Take ordered pair:
 - ① **(0, 0)** neither recently used not modified: **best page** to replace
 - ② **(0, 1)** not recently used but modified: not quite as good, **must write out before replacement**
 - ③ **(1, 0)** recently used but clean: **probably will be used again soon**
 - ④ **(1, 1)** recently used and modified: **probably will be used again soon** and need to write out before replacement
- ▶ When page replacement called for, use the clock scheme but use the four classes replace page in lowest non-empty class

Enhanced Second-Chance Algorithm

- ▶ Improve algorithm by using reference bit and modify bit (reference, modify)
- ▶ Take ordered pair:
 - ① (0, 0) neither recently used nor modified: best page to replace
 - ② (0, 1) not recently used but modified: not quite as good, must write out before replacement
 - ③ (1, 0) recently used but clean: probably will be used again soon
 - ④ (1, 1) recently used and modified: probably will be used again soon and need to write out before replacement
- ▶ When page replacement called for, use the clock scheme but use the four classes replace page in lowest non-empty class
- ▶ Might need to search circular queue several times.

Counting Page Replacement

Counting Page Replacement

- ▶ Keep a **counter** of the **number of references** that have been made to each page.

Counting Page Replacement

- ▶ Keep a **counter** of the **number of references** that have been made to each page.
- ▶ **Least Frequently Used (LFU)** algorithm: replaces page with **smallest count**.

Counting Page Replacement

- ▶ Keep a counter of the number of references that have been made to each page.
- ▶ **Least Frequently Used (LFU)** algorithm: replaces page with **smallest count**.
- ▶ **Most Frequently Used (MFU)** algorithm: based on the argument that the page with the smallest count was **probably just brought in** and has yet to be used.

Summary

- ▶ Partially-loaded programs

Summary

- ▶ Partially-loaded programs
- ▶ Virtual memory: much larger than physical memory

- ▶ Partially-loaded programs
- ▶ Virtual memory: much larger than physical memory
- ▶ Demand paging similar to paging + swapping

- ▶ Partially-loaded programs
- ▶ Virtual memory: much larger than physical memory
- ▶ Demand paging similar to paging + swapping
- ▶ Page fault

- ▶ Partially-loaded programs
- ▶ Virtual memory: much larger than physical memory
- ▶ Demand paging similar to paging + swapping
- ▶ Page fault
- ▶ Page replacement algorithms:
 - FIFO, optimal, LRU, LRU-approximate, counting-based

Questions?

Acknowledgements

Some slides were derived from Avi Silberschatz slides.