

Virtual Memory (Part II)

Amir H. Payberah
amir@sics.se

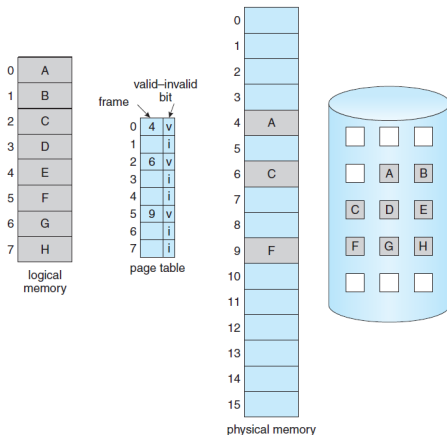
Amirkabir University of Technology
(Tehran Polytechnic)



Reminder

Reminder

- ▶ Partially-loaded programs
- ▶ Virtual memory: much larger than physical memory



- ▶ Demand paging similar to paging + swapping
- ▶ Locality
- ▶ Page fault
- ▶ Page replacement algorithms:
 - FIFO, optimal, LRU, LRU-approximate, counting-based

Allocation of Frames

Allocation of Frames

- ▶ How do we **allocate** the **fixed amount of free memory** among the **various processes**?

Allocation of Frames

- ▶ How do we allocate the fixed amount of free memory among the various processes?
- ▶ If we have 93 free frames and two processes, how many frames does each process get?

Frame Allocation Constraints (1/2)

- ▶ The **maximum** number of frames is the **total frames** in the system.

Frame Allocation Constraints (1/2)

- ▶ The **maximum** number of frames is the **total frames** in the system.
- ▶ Each process needs **minimum** number of **frames**.

Frame Allocation Constraints (1/2)

- ▶ The maximum number of frames is the total frames in the system.
- ▶ Each process needs minimum number of frames.
 - Example: IBM 370: 6 pages to handle MOVE instruction:
 - Instruction is 6 bytes, might span 2 pages
 - 2 pages to handle from
 - 2 pages to handle to

Frame Allocation Constraints (2/2)

- ▶ Why the **minimum** number of frames for each process?

Frame Allocation Constraints (2/2)

- ▶ Why the **minimum** number of frames for each process? **performance**

Frame Allocation Constraints (2/2)

- ▶ Why the **minimum** number of frames for each process? **performance**
- ▶ **Decreases** the number of frames:
 - **Increases** the **page-fault rate**

Frame Allocation Constraints (2/2)

- ▶ Why the **minimum** number of frames for each process? **performance**
- ▶ **Decreases** the number of frames:
 - **Increases** the **page-fault rate**
- ▶ When a **page fault** occurs **before** an executing instruction is **complete**, the instruction must be **restarted**.

Frame Allocation Constraints (2/2)

- ▶ Why the minimum number of frames for each process? performance
- ▶ Decreases the number of frames:
 - Increases the page-fault rate
- ▶ When a page fault occurs before an executing instruction is complete, the instruction must be restarted.
 - We must have enough frames to hold all the different pages that any single instruction can reference.

Allocation Schemes

- ▶ Fixed allocation
- ▶ Priority allocation

Allocation Schemes

- ▶ Fixed allocation
 - Equal allocation
 - Proportional allocation
- ▶ Priority allocation

Fixed Allocation (1/4)

- ▶ Equal allocation
- ▶ Split m frames among n processes: $\frac{m}{n}$ frames to each process.

Fixed Allocation (1/4)

- ▶ Equal allocation
- ▶ Split m frames among n processes: $\frac{m}{n}$ frames to each process.
- ▶ Example, if there are 93 frames and 5 processes
 - Each process will get 18 frames.
 - The 3 leftover frames can be used as a free-frame buffer pool.

Fixed Allocation (2/4)

- ▶ Assume 62 free frames, and the frame size is 1KB

Fixed Allocation (2/4)

- ▶ Assume 62 free frames, and the frame size is 1KB
- ▶ Two processes:
 - A student process: 10KB
 - An interactive database: 127KB

Fixed Allocation (2/4)

- ▶ Assume 62 free frames, and the frame size is 1KB
- ▶ Two processes:
 - A student process: 10KB
 - An interactive database: 127KB
- ▶ Equal allocation: gives each process 31 frames

Fixed Allocation (2/4)

- ▶ Assume 62 free frames, and the frame size is 1KB
- ▶ Two processes:
 - A student process: 10KB
 - An interactive database: 127KB
- ▶ Equal allocation: gives each process 31 frames
- ▶ Wasting 21 frames

Fixed Allocation (3/4)

- ▶ Proportional allocation
- ▶ Allocate according to the size of process.
- ▶ Dynamic as degree of multiprogramming, process sizes change.

Fixed Allocation (3/4)

- ▶ Proportional allocation
- ▶ Allocate according to the size of process.
- ▶ Dynamic as degree of multiprogramming, process sizes change.
- ▶ s_i = size of process p_i

Fixed Allocation (3/4)

- ▶ Proportional allocation
- ▶ Allocate according to the size of process.
- ▶ Dynamic as degree of multiprogramming, process sizes change.
- ▶ s_i = size of process p_i
- ▶ $S = \sum s_i$

Fixed Allocation (3/4)

- ▶ Proportional allocation
- ▶ Allocate according to the size of process.
- ▶ Dynamic as degree of multiprogramming, process sizes change.
- ▶ s_i = size of process p_i
- ▶ $S = \sum s_i$
- ▶ m = total number of frames

Fixed Allocation (3/4)

- ▶ Proportional allocation
- ▶ Allocate according to the size of process.
- ▶ Dynamic as degree of multiprogramming, process sizes change.
- ▶ s_i = size of process p_i
- ▶ $S = \sum s_i$
- ▶ m = total number of frames
- ▶ $a_i = \frac{s_i}{S} \times m$: allocation for p_i

Fixed Allocation (4/4)

- ▶ Assume 62 free frames, and the frame size is 1KB

Fixed Allocation (4/4)

- ▶ Assume 62 free frames, and the frame size is 1KB
- ▶ Two processes:
 - A student process: 10KB
 - An interactive database: 127KB

Fixed Allocation (4/4)

- ▶ Assume 62 free frames, and the frame size is 1KB
- ▶ Two processes:
 - A student process: 10KB
 - An interactive database: 127KB
- ▶ Equal allocation: $s_1 = 10, s_2 = 127, S = 137, m = 62$

Fixed Allocation (4/4)

- ▶ Assume 62 free frames, and the frame size is 1KB
- ▶ Two processes:
 - A student process: 10KB
 - An interactive database: 127KB
- ▶ Equal allocation: $s_1 = 10, s_2 = 127, S = 137, m = 62$
- ▶ $a_1 = \frac{10}{137} \times 62 \approx 4$ $a_2 = \frac{127}{137} \times 62 \approx 57$

Priority Allocation

- ▶ Use a **proportional allocation** scheme using **priorities** rather than **size**.

Priority Allocation

- ▶ Use a **proportional allocation** scheme **using priorities** rather than size.
- ▶ If process p_i generates a **page fault**:
 - Select for replacement one of **its frames**.
 - Select for replacement a frame from a process with **lower priority number**.

Global vs. Local Allocation

- ▶ **Global replacement:** process selects a replacement frame from the set of **all frames**; one process can take a frame from another

Global vs. Local Allocation

- ▶ **Global replacement:** process selects a replacement frame from the set of **all frames**; one process can take a frame from another
 - The process **execution time** can **vary greatly**.
 - The **greater throughput** so more common.

Global vs. Local Allocation

- ▶ **Global replacement:** process selects a replacement frame from the set of **all frames**; one process can take a frame from another
 - The process **execution time** can **vary greatly**.
 - The **greater throughput** so more common.
- ▶ **Local replacement:** each process selects from only **its own set of allocated frames**

Global vs. Local Allocation

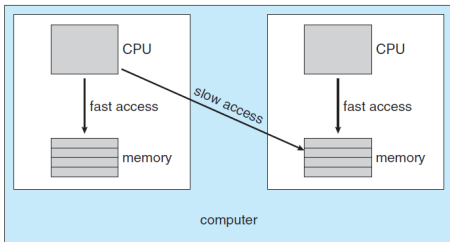
- ▶ **Global replacement:** process selects a replacement frame from the set of all frames; one process can take a frame from another
 - The process execution time can vary greatly.
 - The greater throughput so more common.

?

- ▶ **Local replacement:** each process selects from only its own set of allocated frames
 - More consistent per-process performance
 - Possibly underutilized memory

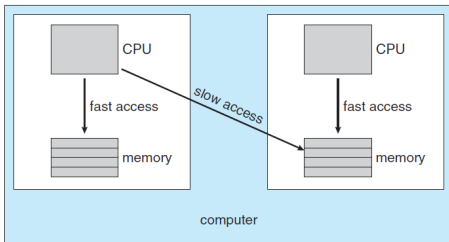
Non-Uniform Memory Access

- ▶ So far all memory accessed **equally**.
- ▶ Many systems are **NUMA**: speed of access to memory **varies**.



Non-Uniform Memory Access

- ▶ So far all memory accessed **equally**.
- ▶ Many systems are **NUMA**: speed of access to memory **varies**.



- ▶ **Optimal performance**: allocate memory **close to** the CPU on which the thread is scheduled.

Thrashing

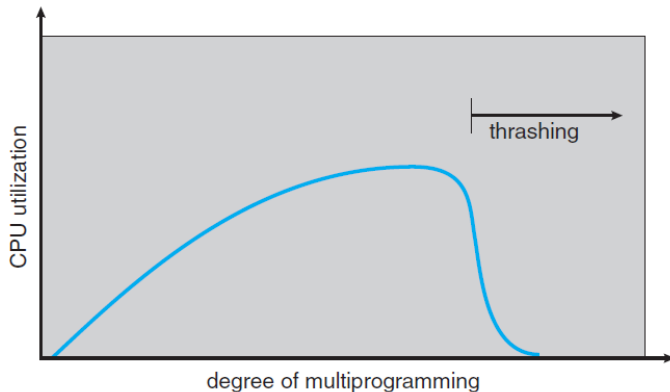
Thrashing (1/2)

- ▶ If a process does not have enough pages, the page-fault rate is very high.
 - Page fault to get page
 - Replace existing frame
 - But quickly need replaced frame back

Thrashing (1/2)

- ▶ If a process does not have enough pages, the page-fault rate is very high.
 - Page fault to get page
 - Replace existing frame
 - But quickly need replaced frame back
- ▶ This leads to:
 - Low CPU utilization
 - OS thinks it needs to increase the degree of multiprogramming
 - Another process added to the system

Thrashing (2/2)



- Thrashing: a process is busy swapping pages in and out.

Prevent Thrashing

- ▶ Providing a process with as many frames as it needs.

Prevent Thrashing

- ▶ Providing a process with as many frames as it needs.
- ▶ How do we know how many frames it needs?

Locality Model (1/2)

- ▶ A **locality** is a **set of pages** that are **actively used together**.

Locality Model (1/2)

- ▶ A **locality** is a **set of pages** that are **actively used together**.
- ▶ A **process** moves from **locality to locality**.

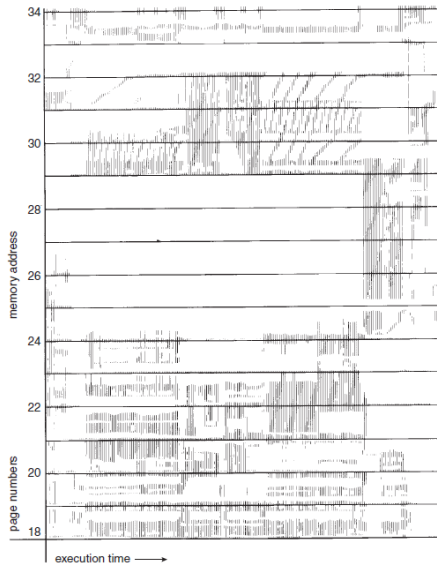
Locality Model (1/2)

- ▶ A **locality** is a **set of pages** that are **actively used together**.
- ▶ A **process** moves from **locality to locality**.
- ▶ A **program** is generally composed of **several different localities**, which may **overlap**.

Locality Model (1/2)

- ▶ A **locality** is a **set of pages** that are **actively used together**.
- ▶ A **process** moves from **locality to locality**.
- ▶ A **program** is generally composed of **several different localities**, which may **overlap**.
- ▶ For example, when a **function** is called, it defines a **new locality**: consists of **memory references** to the instructions of the function call, its **local variables**, and a subset of the **global variables**.

Locality Model (2/2)



Locality and Thrashing

- ▶ A process will fault for the pages in its locality, until all these pages are in memory.

Locality and Thrashing

- ▶ A process will **fault for the pages** in its **locality**, until all these pages are in **memory**.
- ▶ After allocating **all the pages** of the locality, it will **not fault again** until it changes localities.

Locality and Thrashing

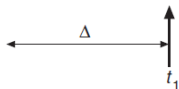
- ▶ A process will fault for the pages in its locality, until all these pages are in memory.
- ▶ After allocating all the pages of the locality, it will not fault again until it changes localities.
- ▶ If we do not allocate enough frames to accommodate the size of the current locality, the process will thrash.

Working-Set Model (1/2)

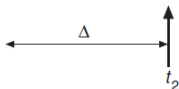
- ▶ Δ : working-set window: a fixed number of page references
- ▶ WSS_i : working set of process p_i : total number of pages referenced in the most recent Δ (varies in time).

page reference table

... 2 6 1 5 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...



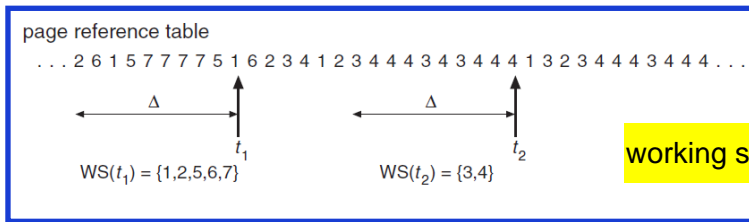
$$WS(t_1) = \{1, 2, 5, 6, 7\}$$



$$WS(t_2) = \{3, 4\}$$

Working-Set Model (1/2)

- ▶ Δ : working-set window: a fixed number of page references
- ▶ WSS_i : working set of process p_i : total number of pages referenced in the most recent Δ (varies in time).
 - If Δ too small will not encompass entire locality
 - If Δ too large will encompass several localities
 - If $\Delta = \infty$ will encompass entire program



working set = 10

Working-Set Model (2/2)

- ▶ m : total number of frames
- ▶ D : total demand frames: $D = \sum WSS_i$:
 - Approximation of locality

Working-Set Model (2/2)

- ▶ m : total number of frames
- ▶ D : total demand frames: $D = \sum WSS_i$:
 - Approximation of locality
- ▶ If $D > m$, then thrashing

Working-Set Model (2/2)

?

- ▶ m : total number of frames
- ▶ D : total demand frames: $D = \sum WSS_i$:
 - Approximation of locality
- ▶ If $D > m$, then thrashing
- ▶ Policy: if $D > m$, then suspend or swap out one of the processes.

Keeping Track of the Working Set

- ▶ Approximate with **interval timer** + a **reference bit**

Keeping Track of the Working Set

- ▶ Approximate with **interval timer** + a **reference bit**
- ▶ Example: $\Delta = 10000$
 - **Timer interrupts** after every **5000** time units.
 - Keep in memory **2 bits** for **each page**.
 - Whenever a timer interrupts we copy and clear the reference-bit values for each page
 - If a page fault occurs: examine the 2 bits to determine whether a page was used within the last 10,000 to 15,000 references.

Keeping Track of the Working Set

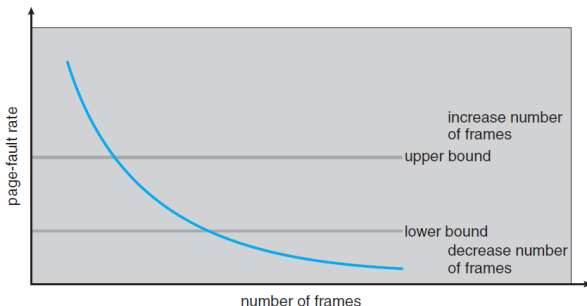
- ▶ Approximate with **interval timer** + a **reference bit**
- ▶ Example: $\Delta = 10000$
 - **Timer interrupts** after every **5000** time units.
 - Keep in memory **2 bits** for **each page**.
 - Whenever a timer interrupts we copy and clear the reference-bit values for each page
 - If a page fault occurs: examine the 2 bits to determine whether a page was used within the last 10,000 to 15,000 references.
- ▶ Why is this not completely accurate?

Keeping Track of the Working Set

- ▶ Approximate with interval timer + a reference bit
- ▶ Example: $\Delta = 10000$
 - Timer interrupts after every 5000 time units.
 - Keep in memory 2 bits for each page.
 - Whenever a timer interrupts we copy and clear the reference-bit values for each page
 - If a page fault occurs: examine the 2 bits to determine whether a page was used within the last 10,000 to 15,000 references.
- ▶ Why is this not completely accurate?
- ▶ Improvement = 10 bits and interrupt every 1000 time units

Page-Fault Frequency

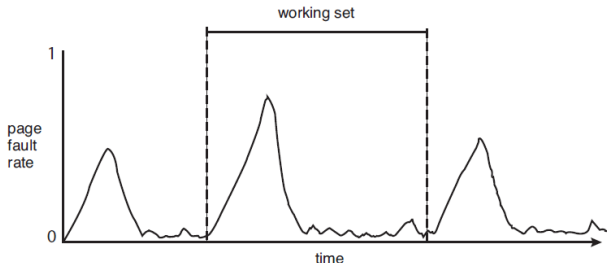
- ▶ More direct approach than WSS
 - ▶ Establish **acceptable page-fault frequency (PFF)** rate and use local replacement policy
- ?
- If **actual rate** too **low**, process loses frame
 - If **actual rate** too **high**, process gains frame



Working Sets and Page-Fault Rates

?

- ▶ Direct relationship between working set of a process and its page-fault rate.
- ▶ Working set changes over time.
- ▶ Peaks and valleys over time.



Allocating Kernel Memory

Allocating Kernel Memory

- ▶ Treated differently from user memory

Allocating Kernel Memory

- ▶ Treated differently from user memory
- ▶ Often allocated from a free-memory pool
 - Kernel requests memory for structures of varying sizes
 - Some kernel memory needs to be contiguous, i.e. for I/O devices

Managing Free Memory Strategies

- ▶ Buddy system
- ▶ Slab allocation

Buddy System (1/2)

- ▶ Allocates memory from **fixed-size segment** consisting of **physically-contiguous pages**

Buddy System (1/2)

- ▶ Allocates memory from **fixed-size segment** consisting of **physically-contiguous pages**
- ▶ Memory allocated using **power-of-2 allocator**

Buddy System (1/2)

- ▶ Allocates memory from **fixed-size segment** consisting of **physically-contiguous pages**
- ▶ Memory allocated using **power-of-2 allocator**
 - Satisfies requests in **units** sized as **power of 2**.

Buddy System (1/2)

- ▶ Allocates memory from **fixed-size segment** consisting of **physically-contiguous pages**
- ▶ Memory allocated using **power-of-2 allocator**
 - Satisfies requests in **units** sized as **power of 2**.
 - Request **rounded up** to **next highest power of 2**.

Buddy System (1/2)

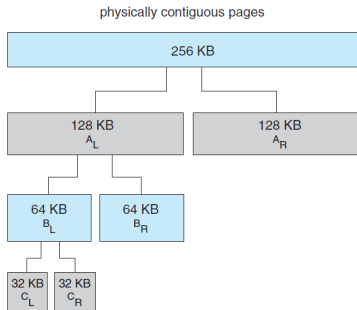
- ▶ Allocates memory from **fixed-size segment** consisting of **physically-contiguous pages**
- ▶ Memory allocated using **power-of-2 allocator**
 - Satisfies requests in **units** sized as **power of 2**.
 - Request **rounded up** to **next highest power of 2**.
 - When **smaller allocation** needed than is available, current chunk **split into two buddies** of next-lower power of 2.

Buddy System (1/2)

- ▶ Allocates memory from fixed-size segment consisting of physically-contiguous pages
- ▶ Memory allocated using power-of-2 allocator
 - Satisfies requests in units sized as power of 2.
 - Request rounded up to next highest power of 2.
 - When smaller allocation needed than is available, current chunk split into two buddies of next-lower power of 2.
 - Continue until appropriate sized chunk available.

Buddy System (2/2)

- ▶ Assume 256KB chunk available, kernel requests 21KB.
 - Split into AL and AR of 128KB each.
 - One further divided into BL and BR of 64KB.
 - One further into CL and CR of 32KB each.

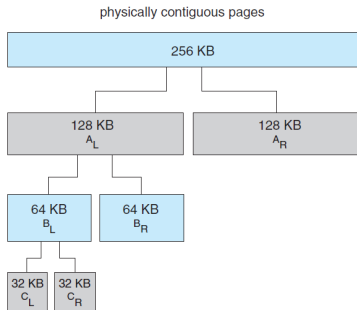


Buddy System (2/2)

- ▶ Assume 256KB chunk available, kernel requests 21KB.
 - Split into AL and AR of 128KB each.
 - One further divided into BL and BR of 64KB.
 - One further into CL and CR of 32KB each.

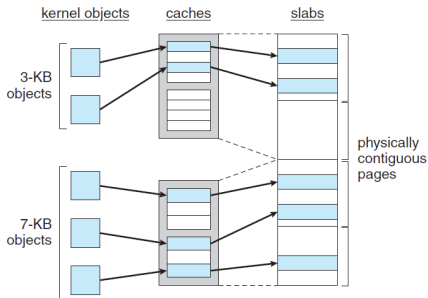
? ▶ Advantage: quickly coalesce unused chunks into larger chunk

- ▶ Disadvantage: fragmentation



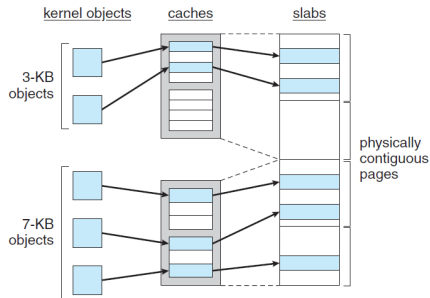
Slab Allocation (1/2)

- ▶ Alternate strategy
- ▶ **Slab** is one or more **physically contiguous pages**.



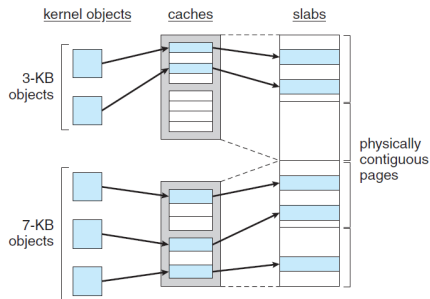
Slab Allocation (1/2)

- ▶ Alternate strategy
- ▶ **Slab** is one or more **physically contiguous pages**.
- ▶ **Cache** consists of **one or more slabs**.



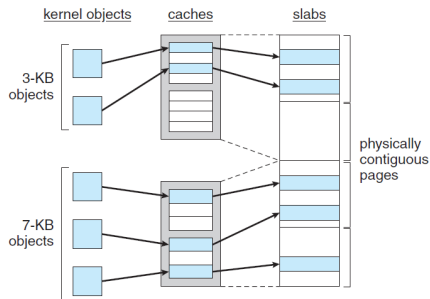
Slab Allocation (1/2)

- ▶ Alternate strategy
- ▶ **Slab** is one or more **physically contiguous pages**.
- ▶ **Cache** consists of **one or more slabs**.
- ▶ **Single cache** for each unique **kernel data structure**, e.g., a separate cache for file objects, a separate cache for semaphores, and so forth.



Slab Allocation (1/2)

- ▶ Alternate strategy
- ▶ Slab is one or more physically contiguous pages.
- ▶ Cache consists of one or more slabs.
- ▶ Single cache for each unique kernel data structure, e.g.,
a separate cache for file objects,
a separate cache for semaphores,
and so forth.
- ▶ Objects: instantiations of the data structure



Slab Allocation (2/2)

- ▶ When `cache` created, filled with objects marked as `free`.

Slab Allocation (2/2)

- ▶ When **cache** created, filled with objects marked as **free**.
- ▶ When **structures stored**, objects marked as **used**.

Slab Allocation (2/2)

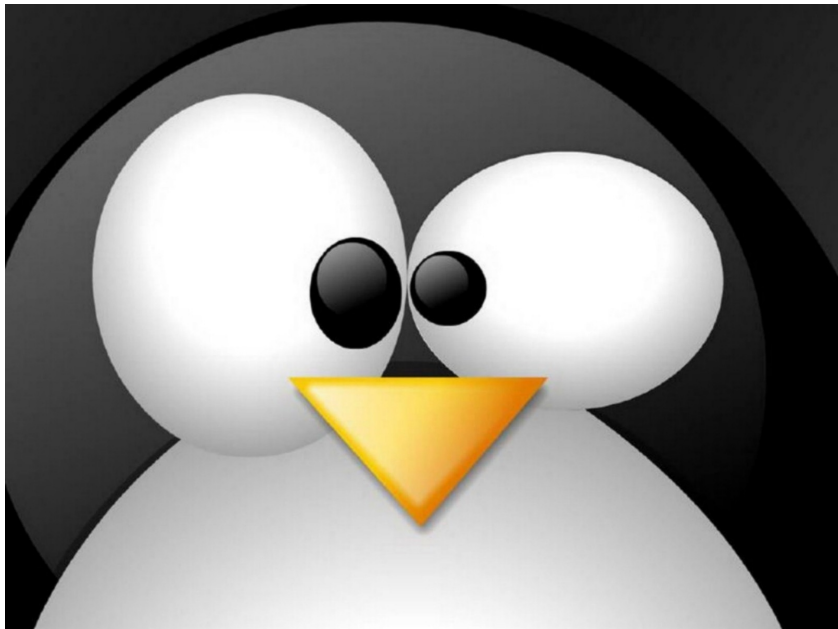
- ▶ When **cache** created, filled with objects marked as **free**.
- ▶ When **structures stored**, objects marked as **used**.
- ▶ If **slab** is **full** of used objects, next object allocated from **empty slab**.

Slab Allocation (2/2)

- ▶ When **cache** created, filled with objects marked as **free**.
- ▶ When **structures stored**, objects marked as **used**.
- ▶ If **slab** is **full** of used objects, next object allocated from **empty slab**.
- ▶ If no empty slabs, **new slab allocated**.

Slab Allocation (2/2)

- ▶ When **cache** created, filled with objects marked as **free**.
- ▶ When **structures stored**, objects marked as **used**.
- ▶ If **slab** is **full** of used objects, next object allocated from **empty slab**.
- ▶ If no empty slabs, **new slab allocated**.
- ▶ Benefits include **no fragmentation** and **fast** memory request satisfaction.



Slab Allocator in Linux (1/3)

- ▶ Process descriptor: `type struct task_struct`
- ▶ Approx 1.7KB of memory
- ▶ New task → allocate new struct from `cache`
 - Will use existing free `struct task_struct`
- ▶ `Slab` can be in `three` possible states
 - `Full`: all used
 - `Empty`: all free
 - `Partial`: mix of free and used

Slab Allocator in Linux (2/3)

- ▶ Upon request, slab allocator
 - ① Uses free struct in **partial slab**.
 - ② If none, takes one from **empty slab**.
 - ③ If no empty slab, create **new** empty.

Slab Allocator in Linux (3/3)

- ▶ Linux originally used the **buddy** system.

Slab Allocator in Linux (3/3)

- ▶ Linux originally used the **buddy** system.
- ▶ **Kernel 2.2** had **SLAB**.

Slab Allocator in Linux (3/3)

- ▶ Linux originally used the **buddy** system.
- ▶ **Kernel 2.2** had **SLAB**.
- ▶ Recent distribution added **two more allocators**:
 - **SLOB** (Simple List of Blocks): for systems with **limited memory**, maintains **3 list objects** for small, medium, large objects
 - **SLUB** is **performance-optimized SLAB**, removes per-CPU queues, metadata stored in page structure, from **kernel 2.6.24**

Memory-Mapped Files

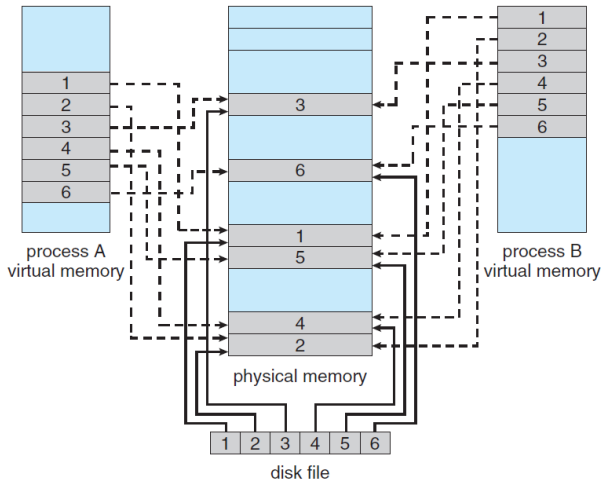
Memory-Mapped Files (1/3)

- ▶ Memory-mapped file I/O allows file I/O to be treated as routine memory access by mapping a disk block to a page in memory.

Memory-Mapped Files (1/3)

- ▶ Memory-mapped file I/O allows file I/O to be treated as routine memory access by mapping a disk block to a page in memory.
- ▶ A file is initially read using demand paging.
 - A page-sized portion of the file is read from the file system into a physical page.
 - Subsequent reads/writes to/from the file are treated as ordinary memory accesses.

Memory-Mapped Files (2/3)



Memory-Mapped Files (3/3)

- ▶ Process can **explicitly request memory mapping** a file via `mmap()` system call: map the file into **kernel address space**

Memory-Mapped Files (3/3)

- ▶ Process can **explicitly request memory mapping** a file via `mmap()` system call: map the file into **kernel address space**
- ▶ **Simplifies** and **speeds file access** by driving file I/O through **memory** rather than `read()` and `write()` system calls.

Memory-Mapped Files (3/3)

- ▶ Process can **explicitly request memory mapping** a file via `mmap()` system call: map the file into **kernel address space**
- ▶ **Simplifies** and **speeds file access** by driving file I/O through **memory** rather than `read()` and `write()` system calls.
- ▶ Also allows **several processes** to map the same file allowing the pages in memory to be **shared**.

Memory-Mapped Files (3/3)

- ▶ Process can **explicitly request memory mapping** a file via `mmap()` system call: map the file into **kernel address space**
- ▶ **Simplifies** and **speeds file access** by driving file I/O through **memory** rather than `read()` and `write()` system calls.
- ▶ Also allows **several processes** to map the same file allowing the pages in memory to be **shared**.
- ▶ **When** does written data make it to disk?

Memory-Mapped Files (3/3)

- ▶ Process can explicitly request memory mapping a file via mmap() system call: map the file into kernel address space
- ▶ Simplifies and speeds file access by driving file I/O through memory rather than read() and write() system calls.
- ▶ Also allows several processes to map the same file allowing the pages in memory to be shared.
- ▶ When does written data make it to disk?
 - Periodically and/or at file close() time.

- ▶ Many computer architectures provide **memory-mapped I/O**.

Memory-Mapped I/O

- ▶ Many computer architectures provide **memory-mapped I/O**.
- ▶ Ranges of **memory addresses** are mapped to the **device registers**.

Memory-Mapped I/O

- ▶ Many computer architectures provide **memory-mapped I/O**.
- ▶ Ranges of **memory addresses** are mapped to the **device registers**.
- ▶ Reads and writes to these memory addresses cause the data to be transferred to and from the device registers.
 - Called, **I/O port**

Other Considerations

- ▶ To **reduce** the large number of **page faults** that occurs at **process startup**.

Prepaging

- ▶ To **reduce** the large number of **page faults** that occurs at **process startup**.
- ▶ **Prepage** all or some of the pages a process will need, **before** they are referenced.

Prepaging

- ▶ To **reduce** the large number of **page faults** that occurs at **process startup**.
- ▶ **Prepage** all or some of the pages a process will need, **before** they are referenced.
- ▶ If prepaged pages are **unused**, I/O and memory was **wasted**.

Prepaging

- ▶ To reduce the large number of page faults that occurs at process startup.
- ▶ Prepage all or some of the pages a process will need, before they are referenced.
- ▶ If prepaged pages are unused, I/O and memory was wasted.
- ▶ Assume s pages are prepaged and a fraction $0 \leq \alpha \leq 1$ of the pages are used.
 - Cost of $s \times \alpha$ > or < than the cost of prepaging $s \times (1 - \alpha)$ unnecessary pages?
 - If α close to 0: prepaging loses; if α close to 1, prepaging wins

- ▶ Page size selection must take into consideration:
 - Fragmentation
 - Page table size
 - Resolution
 - I/O overhead
 - Number of page faults
 - Locality
 - TLB size and effectiveness
- ▶ Always power of 2

- ▶ **TLB reach**: the amount of memory accessible from the TLB

- ▶ **TLB reach**: the amount of memory accessible from the TLB
- ▶ $\text{TLB reach} = (\text{TLB Size}) \times (\text{Page Size})$

- ▶ **TLB reach**: the amount of memory accessible from the TLB
- ▶ $\text{TLB reach} = (\text{TLB Size}) \times (\text{Page Size})$
- ▶ Ideally, the **working set of each process** is stored in the TLB
 - Otherwise there is a high degree of **page faults**.

- ▶ **TLB reach**: the amount of memory accessible from the TLB
- ▶ $\text{TLB reach} = (\text{TLB Size}) \times (\text{Page Size})$
- ▶ Ideally, the **working set of each process** is stored in the TLB
 - Otherwise there is a high degree of **page faults**.
- ▶ Increase the **page size**
 - This may lead to an **increase in fragmentation**

TLB Search

- ▶ **TLB reach**: the amount of memory accessible from the TLB
- ▶ $TLB\ reach = (TLB\ Size) \times (Page\ Size)$
- ▶ Ideally, the **working set of each process** is stored in the TLB
 - Otherwise there is a high degree of **page faults**.
- ▶ Increase the page size
 - This may lead to an increase in fragmentation
- ▶ Provide **multiple page sizes**
 - This allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation

Program Structure

- ▶ `int[128,128] data`: each row is stored in one page
- ▶ Program 1

```
for (j = 0; j < 128; j++)  
    for (i = 0; i < 128; i++)  
        data[i, j] = 0;
```

Program Structure

- ▶ `int[128,128] data`: each row is stored in one page
- ▶ Program 1

```
for (j = 0; j < 128; j++)  
    for (i = 0; i < 128; i++)  
        data[i, j] = 0;
```

$128 \times 128 = 16,384$ page faults

Program Structure

- ▶ `int[128,128] data`: each row is stored in one page
- ▶ Program 1

```
for (j = 0; j < 128; j++)  
    for (i = 0; i < 128; i++)  
        data[i, j] = 0;
```

$128 \times 128 = 16,384$ page faults

- ▶ Program 2

```
for (i = 0; i < 128; i++)  
    for (j = 0; j < 128; j++)  
        data[i, j] = 0;
```

Program Structure

- ▶ `int[128,128] data`: each row is stored in one page
- ▶ Program 1

```
for (j = 0; j < 128; j++)  
    for (i = 0; i < 128; i++)  
        data[i, j] = 0;
```

$128 \times 128 = 16,384$ page faults

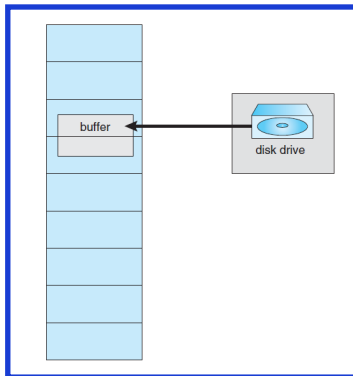
- ▶ Program 2

```
for (i = 0; i < 128; i++)  
    for (j = 0; j < 128; j++)  
        data[i, j] = 0;
```

128 page faults

I/O Interlock

- ▶ **I/O interlock:** pages must sometimes be locked into memory.
- ▶ **Consider I/O:** pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm.



Summary

- ▶ Frame allocation: fixed and priority allocations

- ▶ Frame allocation: fixed and priority allocations
- ▶ Global and local allocation

Summary

- ▶ Frame allocation: fixed and priority allocations
- ▶ Global and local allocation
- ▶ Thrashing: total demand frames $>$ total num. of frames

- ▶ Frame allocation: fixed and priority allocations
- ▶ Global and local allocation
- ▶ Thrashing: total demand frames $>$ total num. of frames
- ▶ Prevent trashing: working set model and page fault frequency

- ▶ Frame allocation: fixed and priority allocations
- ▶ Global and local allocation
- ▶ Thrashing: total demand frames $>$ total num. of frames
- ▶ Prevent trashing: working set model and page fault frequency
- ▶ Allocating kernel memory: buddy system and slab allocation

- ▶ Frame allocation: fixed and priority allocations
- ▶ Global and local allocation
- ▶ Thrashing: total demand frames $>$ total num. of frames
- ▶ Prevent trashing: working set model and page fault frequency
- ▶ Allocating kernel memory: buddy system and slab allocation
- ▶ Memory-mapped files and I/O

Questions?

Acknowledgements

Some slides were derived from Avi Silberschatz slides.