

# File System Implementation (Part I)

Amir H. Payberah  
amir@sics.se

Amirkabir University of Technology  
(Tehran Polytechnic)



# Motivation

# Motivation

- ▶ The file system resides permanently on secondary storage.
- ▶ How to
  - structure file use
  - allocate disk space
  - recover free space
  - track the locations of data
  - interface other parts of the OS to secondary storage

# File System Structure

# File-System Structure

- ▶ **Disk** provides **in-place** rewrite and **random** access
  - I/O transfers performed in **blocks of sectors** (usually 512 bytes)

# File-System Structure

- ▶ **Disk** provides **in-place** rewrite and **random** access
  - I/O transfers performed in **blocks of sectors** (usually 512 bytes)
- ▶ **File system** resides on secondary storage
  - User **interface** to storage, mapping **logical to physical**
  - **Efficient and convenient** access to disk

# File-System Structure

- ▶ **Disk** provides **in-place** rewrite and **random** access
  - I/O transfers performed in **blocks of sectors** (usually 512 bytes)
- ▶ **File system** resides on secondary storage
  - **User interface** to storage, mapping **logical to physical**
  - **Efficient and convenient** access to disk
- ▶ **File structure**
  - **Logical storage unit**
  - Collection of **related information**

# File-System Design Problems

- ▶ How the file system should look to the user?



# File-System Design Problems


- ▶ How the **file system** should **look to the user**?
  - Defining a **file** and its **attributes**
  - The **operations** allowed on a file
  - The **directory structure** for organizing files

# File-System Design Problems

- ▶ How the **file system** should **look to the user**?
  - Defining a file and its attributes
  - The operations allowed on a file
  - The directory structure for organizing files
- ▶ Algorithms and data structures to **map the logical file system onto the physical secondary-storage** devices.


# File System Layers (1/6)

- ▶ Different **levels**
- ▶ Each level uses the features of lower levels to create new features for use by higher levels.
- ▶ Reducing complexity and redundancy, but **adds overhead** and can **decrease performance**.




## File System Layers (2/6)

- ▶ Device drivers manage I/O devices at the I/O control layer.
- ▶ Translates high-level commands to low-level hardware-specific instructions.




## File System Layers (3/6)

- **Basic file system** translates given command like **retrieve block 123** to device driver.




# File System Layers (3/6)

- ▶ **Basic file system** translates given command like **retrieve block 123** to **device driver**.
- ▶ Also **manages memory buffers and caches** (allocation, freeing, replacement)
  - **Buffers** hold data in transit
  - **Caches** hold frequently used data




## File System Layers (4/6)

- **File organization** understands files, logical address, and physical blocks.



## File System Layers (4/6)


- ▶ **File organization** understands files, logical address, and physical blocks.
- ▶ **Translates** logical block number to physical block number.






# File System Layers (4/6)

- ▶ File organization understands files, logical address, and physical blocks.
- ▶ Translates logical block number to physical block number.
- ▶ Manages free space and disk allocation.




# File System Layers (5/6)

- Logical file system manages metadata information.




# File System Layers (5/6)

- ▶ Logical file system manages metadata information.
- ▶ Translates file name into file number, file handle, location by maintaining file control blocks (inodes in Unix)



# File System Layers (5/6)

- ▶ Logical file system manages metadata information.
- ▶ Translates file name into file number, file handle, location by maintaining file control blocks (inodes in Unix)
- ▶ Directory management
- ▶ Protection



# File System Layers (6/6)

- ▶ Many file systems, sometimes many within an OS
- ▶ Each with its own **format**
  - **CD-ROM**: ISO 9660
  - **Unix**: UFS, FFS
  - **Windows**: FAT, FAT32, NTFS
  - **Linux**: more than 40 types, with **extended file system** (ext2, ext3, ext4)

# File System Implementation

# File-System Implementation

- ▶ Based on several **on-disk** and **in-memory** structures.

# File-System Implementation

- ▶ Based on several **on-disk** and **in-memory** structures.
- ▶ On-disk
  - Boot control block (per volume)
  - Volume control block (per volume)
  - Directory structure (per file system)
  - File control block (per file)




# File-System Implementation

- ▶ Based on several on-disk and in-memory structures.
- ▶ On-disk
  - Boot control block (per volume)
  - Volume control block (per volume)
  - Directory structure (per file system)
  - File control block (per file)
- ▶ In-memory
  - Mount table
  - Directory structure cache
  - The open-file table (system-wide and per process)
  - Buffers of the file-system blocks


# On-Disk File System Structures (1/2)

- ▶ **Boot control block** contains information needed by system to **boot OS** from that volume.




# On-Disk File System Structures (1/2)

- ▶ **Boot control block** contains information needed by system to **boot OS** from that volume.
  - Needed if volume **contains OS**, usually **first block of volume**.
  - In **UFS**, it is called **boot block**, and in **NTFS** **partition boot sector**.




# On-Disk File System Structures (1/2)

- ▶ **Boot control block** contains information needed by system to **boot OS** from that volume.
  - Needed if volume **contains OS**, usually **first block of volume**.
  - In **UFS**, it is called **boot block**, and in **NTFS** **partition boot sector**.
- ▶ **Volume control block** contains volume details.



# On-Disk File System Structures (1/2)

- ▶ **Boot control block** contains information needed by system to boot OS from that volume.
  - Needed if volume contains OS, usually first block of volume.
  - In **UFS**, it is called boot block, and in **NTFS** partition boot sector.
- ▶ **Volume control block** contains volume details.
  - Total num. of blocks, num. of free blocks, block size, free block pointers or array
  - In **UFS**, it is called super block, and in **NTFS** master file table.



## On-Disk File System Structures (2/2)

- ▶ **Directory structure** organizes the files.
  - In **UFS**, this includes **file names** and associated **inode numbers**.
  - In **NTFS**, it is stored in the **master file table**.

# On-Disk File System Structures (2/2)

- ▶ **Directory structure** organizes the files.
  - In **UFS**, this includes **file names** and associated **inode numbers**.
  - In **NTFS**, it is stored in the **master file table**.
- ▶ **File control block** contains many details about the file.
  - In **UFS**, inode number, permissions, size, dates.
  - In **NFTS** stores into in **master file table**.

file permissions
file dates (create, access, write)
file owner, group, ACL
file size
file data blocks or pointers to file data blocks

# In-Memory File System Structures

- ▶ **Mount table** contains information about each **mounted volume**.



# In-Memory File System Structures

- ▶ **Mount table** contains information about each **mounted volume**.
- ▶ **Directory structure cache** holds the **directory information** of **recently** accessed directories.

# In-Memory File System Structures

- ▶ **Mount table** contains information about each **mounted volume**.
- ▶ **Directory structure cache** holds the **directory information** of **recently** accessed directories.
- ▶ **System-wide open-file table** contains a copy of the **FCB** of each open file.

# In-Memory File System Structures

- ▶ **Mount table** contains information about each **mounted volume**.
- ▶ **Directory structure cache** holds the **directory information** of **recently** accessed directories.
- ▶ **System-wide open-file table** contains a copy of the **FCB** of each open file.
- ▶ **Per-process open-file table** contains a pointer to the appropriate entry in the **system-wide open-file table**.

# In-Memory File System Structures

- ▶ **Mount table** contains information about each mounted volume.
- ▶ **Directory structure cache** holds the **directory information** of **recently accessed directories**.
- ▶ **System-wide open-file table** contains a copy of the **FCB of each open file**.
- ▶ **Per-process open-file table** contains a pointer to the appropriate entry in the system-wide open-file table.
- ▶ **Buffers** hold file-system blocks when they are being read from disk or written to disk.

# Create a File

- ▶ A program calls the **logical file system**.

# Create a File


- ▶ A program calls the **logical file system**.
- ▶ The **logical file system** knows the **format of the directory structures**, and **allocates a new FCB**.

## Create a File

- ▶ A program calls the logical file system.
- ▶ The logical file system knows the format of the directory structures, and allocates a new FCB.
- ▶ The system, then, reads the appropriate directory into memory, updates it with the new file name and FCB, and writes it back to the disk.

# Open a File


- ▶ The file must be **opened**.
  - The **open()** passes a **file name** to the **logical file system**.






# Open a File

- ▶ The file must be **opened**.
  - The **open()** passes a **file name** to the **logical file system**.
- ▶ The **open()** first searches the **system-wide open-file**: if the file is already in use by **another process**.




# Open a File

- ▶ The file must be **opened**.
  - The **open()** passes a **file name** to the **logical file system**.
- ▶ The **open()** first searches the **system-wide open-file**: if the file is already in use by **another process**.
  - If **yes**: a **per-process open-file table** entry is created.




# Open a File

- ▶ The file must be **opened**.
  - The **open()** passes a **file name** to the **logical file system**.
- ▶ The **open()** first searches the **system-wide open-file**: if the file is already in use by **another process**.
  - If **yes**: a **per-process open-file table** entry is created.
  - If **no**: the **directory structure** is searched for the given **file name**: once the file is found, the **FCB** is copied into a **system-wide open-file table in memory**.




# Open a File

- ▶ The file must be **opened**.
  - The **open()** passes a **file name** to the **logical file system**.
- ▶ The **open()** first searches the **system-wide open-file**: if the file is already in use by **another process**.
  - If **yes**: a **per-process open-file table entry** is created.
  - If **no**: the **directory structure** is **searched** for the given **file name**: once the file is found, the **FCB** is **copied** into a **system-wide open-file table in memory**.
- ▶ This table stores the **FCB** as well as the **number of processes** that have the file open.



# Read From a File

- ▶ The open() returns a pointer to the appropriate entry in the per-process file-system table.
- ▶ All file operations are then performed via this pointer.
- ▶ This pointer is called file descriptor in Unix and file handle in Windows.



# Close a File

- ▶ When a process **closes** the file:
  - The **per-process table** entry is **removed**.
  - The **system-wide** entry's open count is **decremented**.

# Close a File

- ▶ When a process **closes** the file:
  - The **per-process table** entry is **removed**.
  - The **system-wide** entry's open count is **decremented**.
- ▶ When **all users** that have opened the file **close** it, **any updated meta-data** is **copied back to the disk-based directory structure**, and the **system-wide** open-file table entry is **removed**.

# Partitions and Mounting (1/2)

- ▶ **Partition** can be a **volume** containing a **file system** or **raw**.
  - **Raw** partition: just a **sequence of blocks** with no file system.



# Partitions and Mounting (1/2)

- ▶ **Partition** can be a **volume** containing a file system or raw.
  - Raw partition: just a sequence of blocks with no file system.
- ▶ **Boot block** points to boot volume or boot loader.
  - **Boot loader:** knows enough about the file-system structure to be able to find and load the kernel and start it executing.
  - **Dual-boot** that allows to install multiple OS on a single system.

## Partitions and Mounting (2/2)

- ▶ **Root partition** contains the OS
  - Mounted at boot time
  - Other partitions can hold other OSes, other file systems, or be raw
  - Other partitions can mount automatically or manually

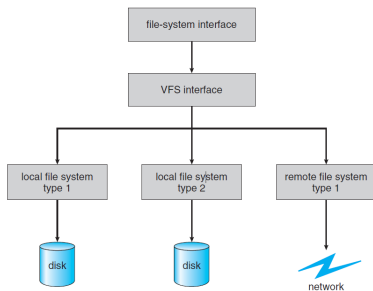
# Partitions and Mounting (2/2)

- ▶ Root partition contains the OS
  - Mounted at boot time
  - Other partitions can hold other OSes, other file systems, or be raw
  - Other partitions can mount automatically or manually
- ▶ At mount time, file system consistency checked.
  - Is all metadata correct? if not, fix it, try again, if yes, add to mount table, allow access

# Virtual File Systems


# Virtual File Systems (1/2)

- **Virtual File Systems (VFS)** on Unix provide an **object-oriented** way of **implementing** file systems.



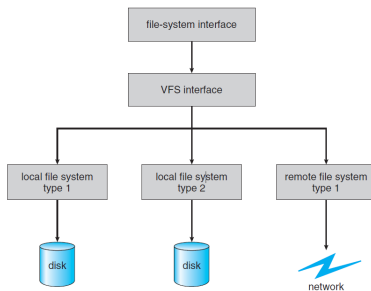
# Virtual File Systems (1/2)

- ▶ **Virtual File Systems (VFS)** on Unix provide an **object-oriented** way of **implementing** file systems.
- ▶ **VFS** allows the **same system call interface** (the API) to be used for **different types of file systems**.



# Virtual File Systems (1/2)

- ▶ Virtual File Systems (VFS) on Unix provide an object-oriented way of implementing file systems.
- ▶ VFS allows the same system call interface (the API) to be used for different types of file systems.
- ▶ The API is to the VFS interface, rather than any specific type of file system.



## Virtual File Systems (2/2)

- ▶ VFS layer serves two important functions:



## Virtual File Systems (2/2)

- ▶ VFS layer serves two important functions:
  - ① It separates file-system-generic operations from their implementation, and allows transparent access to different types of file systems mounted locally.

# Virtual File Systems (2/2)

- ▶ VFS layer serves two important functions:
  - ① It separates file-system-generic operations from their implementation, and allows transparent access to different types of file systems mounted locally.
  - ② It provides a mechanism for uniquely representing a file throughout a network.

## Virtual File Systems (2/2)

- ▶ VFS layer serves two important functions:
  - ① It separates file-system-generic operations from their implementation, and allows transparent access to different types of file systems mounted locally.
  - ② It provides a mechanism for uniquely representing a file throughout a network.
- ▶ The VFS is based on a structure, called a vnode.

## Virtual File Systems (2/2)

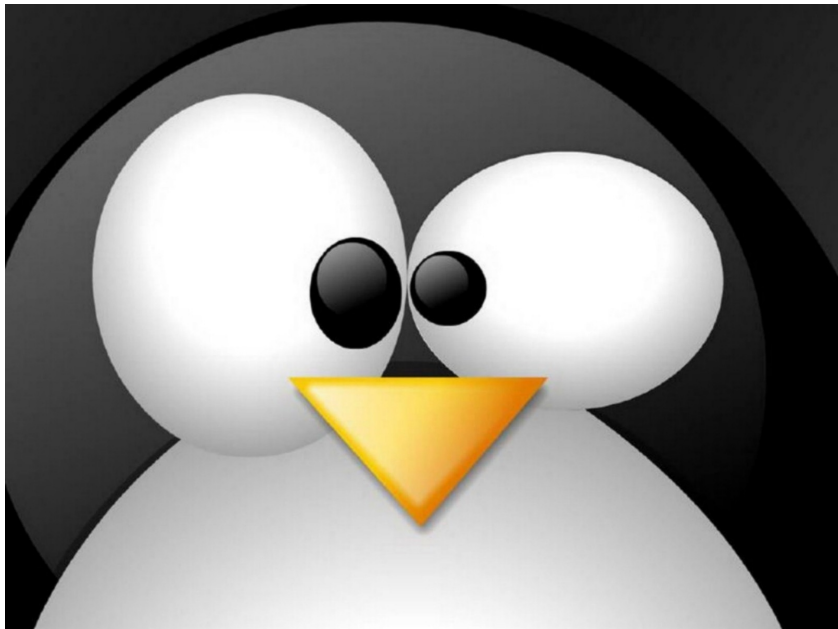
- ▶ VFS layer serves two important functions:
  - ① It separates file-system-generic operations from their implementation, and allows transparent access to different types of file systems mounted locally.
  - ② It provides a mechanism for uniquely representing a file throughout a network.
- ▶ The VFS is based on a structure, called a vnode.
  - Contains a numerical designator for a network-wide unique file.

## Virtual File Systems (2/2)

- ▶ VFS layer serves two important functions:
  - ① It separates file-system-generic operations from their implementation, and allows transparent access to different types of file systems mounted locally.
  - ② It provides a mechanism for uniquely representing a file throughout a network.
- ▶ The VFS is based on a structure, called a vnode.
  - Contains a numerical designator for a network-wide unique file.
  - Unix inodes are unique within only a single file system.

# Virtual File Systems (2/2)

- ▶ VFS layer serves **two important functions**:
  - ① It **separates** **file-system-generic operations** from their **implementation**, and allows **transparent access** to **different types** of file systems mounted locally.
  - ② It provides a **mechanism for uniquely representing a file throughout a network**.
- ▶ The VFS is based on a **structure**, called a **vnode**.
  - Contains a numerical designator for a **network-wide unique file**.
  - Unix **inodes** are unique within only a **single file system**.
  - The kernel maintains **one vnode** structure for **each active node**.



- ▶ The **four** main **object** types defined by the **Linux VFS** are:



- ▶ The **four** main **object** types defined by the **Linux VFS** are:
  - The **inode object**: represents an **individual file**

- ▶ The **four** main **object** types defined by the **Linux VFS** are:
  - The **inode object**: represents an **individual file**
  - The **file** object: represents an **open file**

- ▶ The **four** main **object** types defined by the **Linux VFS** are:
  - The **inode object**: represents an **individual file**
  - The **file** object: represents an **open file**
  - The **super block** object: represents an **entire file system**

- ▶ The **four** main **object** types defined by the **Linux VFS** are:
  - The **inode object**: represents an **individual file**
  - The **file** object: represents an **open file**
  - The **super block** object: represents an **entire file system**
  - The **dentry** object: represents an **individual directory entry**

- ▶ VFS defines a **set of operations** on the **objects** that must be implemented.

## VFS in Linux (2/2)

- ▶ VFS defines a **set of operations** on the **objects** that must be implemented.
- ▶ Every **object** has a **pointer** to a **function table**.

## VFS in Linux (2/2)

- ▶ VFS defines a **set of operations** on the **objects** that must be implemented.
- ▶ Every **object** has a **pointer** to a **function table**.
  - **Function table** has **addresses of routines** to implement that function on that object.

## VFS in Linux (2/2)

- ▶ VFS defines a **set of operations** on the **objects** that must be implemented.
- ▶ Every **object** has a **pointer** to a **function table**.
  - **Function table** has **addresses of routines** to implement that function on that object.
  - For example:
    - `int open(...)`: open a file
    - `int close(...)`: close an already-open file
    - `ssize_t read(...)`: read from a file
    - `ssize_t write(...)`: write to a file
    - `int mmap(...)`: memory-map a file



# Directory Implementation

# Directory Implementation

- ▶ Linear list
- ▶ Hash table

# Directory Implementation - Linear List

- ▶ Linear list of file names with pointer to the data blocks.

# Directory Implementation - Linear List

- ▶ Linear list of file names with pointer to the data blocks.
- ▶ Simple to program.

# Directory Implementation - Linear List

- ▶ Linear list of file names with pointer to the data blocks.
- ▶ Simple to program.
- ▶ Time-consuming to execute.

# Directory Implementation - Linear List

- ▶ Linear list of file names with pointer to the data blocks.
- ▶ Simple to program.
- ▶ Time-consuming to execute.
- ▶ Linear search time.

## Directory Implementation - Linear List

- ▶ Linear list of file names with pointer to the data blocks.
- ▶ Simple to program.
- ▶ Time-consuming to execute.
- ▶ Linear search time.
- ▶ Could keep ordered alphabetically via linked list or use B+ tree: binary search, but heavy

# Directory Implementation - Hash Table

- ▶ **Hash Table:** linear list with hash data structure



# Directory Implementation - Hash Table

- ▶ **Hash Table:** linear list with hash data structure
- ▶ Decreases directory search time

# Directory Implementation - Hash Table

- ▶ **Hash Table**: linear list with hash data structure
- ▶ Decreases directory search time
- ▶ **Collisions**: situations where two file names hash to the same location

# Directory Implementation - Hash Table

- ▶ **Hash Table**: linear list with hash data structure
- ▶ Decreases directory search time
- ▶ **Collisions**: situations where two file names hash to the same location
- ▶ **Chained-overflow method**.
  - Each hash entry can be a **linked list** instead of an individual value.

# Allocation Methods

- ▶ How disk blocks are allocated to files?

▶ How disk blocks are allocated to files?

▶ Methods:

- Contiguous allocation
- Linked allocation
- Indexed allocation

# Contiguous Allocation

# Contiguous Allocation (1/2)

- ▶ **Contiguous allocation:** each file occupies set of contiguous blocks.



# Contiguous Allocation (1/2)

- ▶ **Contiguous allocation:** each file occupies set of contiguous blocks.
  - **Best performance** in most cases
  - **Simple:** only **starting location** (block number) and **length** (number of blocks) are required.


# Contiguous Allocation (1/2)

- ▶ **Contiguous allocation:** each file occupies set of contiguous blocks.
  - Best performance in most cases
  - Simple: only starting location (block number) and length (number of blocks) are required.
  - Supports both sequential and direct access.

# Contiguous Allocation (1/2)

- ▶ **Contiguous allocation:** each file occupies set of contiguous blocks.
  - **Best performance** in most cases
  - **Simple:** only starting location (block number) and length (number of blocks) are required.
  - **Supports** both **sequential** and **direct access**.
- ▶ Allocation strategies like contiguous memory allocation:
  - First fit
  - Best fit
  - Worst fit

## Contiguous Allocation (2/2)



directory

file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2

# Contiguous Allocation Problems

- ▶ Finding space for file

# Contiguous Allocation Problems

- ▶ Finding space for file
- ▶ External fragmentation

# Contiguous Allocation Problems

- ▶ Finding space for file
- ▶ External fragmentation
- ▶ Need for compaction (fragmentation) off-line or on-line: lose of performance

# Contiguous Allocation Problems

- ▶ Finding space for file
- ▶ External fragmentation
- ▶ Need for compaction (fragmentation) off-line or on-line: lose of performance
- ? ▶ Knowing file size



# Extent-Based Systems

- ▶ A modified contiguous allocation scheme.
  - E.g., Veritas file system

# Extent-Based Systems

- ▶ A **modified contiguous allocation** scheme.
  - E.g., Veritas file system
- ▶ **Extent-based** file systems allocate disk blocks in **extents**.

# Extent-Based Systems

- ▶ A modified contiguous allocation scheme.
  - E.g., Veritas file system
- ▶ Extent-based file systems allocate disk blocks in extents.
- ▶ An extent is a contiguous block of disks.
  - Extents are allocated for file allocation.
  - A file consists of one or more extents.

# Linked Allocation

# Linked Allocation (1/2)

- ▶ **Linked allocation:** each file is a **linked list of blocks**.
  - Each block contains **pointer to next block**.
  - File ends at **null pointer**.


# Linked Allocation (1/2)

- ▶ **Linked allocation**: each file is a **linked list of blocks**.
  - Each block contains **pointer to next block**.
  - File ends at **null pointer**.
- ▶ **No external fragmentation**, **no compaction**.

# Linked Allocation (1/2)

- ▶ **Linked allocation:** each file is a linked list of blocks.
  - Each block contains pointer to next block.
  - File ends at null pointer.
- ▶ No external fragmentation, no compaction.
- ▶ Free space management system called when new block needed.

## Linked Allocation (2/2)





# Linked Allocation Problems

- ▶ Locating a block can take many I/Os and disk seeks.

# Linked Allocation Problems


- ▶ Locating a block can take many I/Os and disk seeks.
- ▶ Reliability can be a problem.

# Linked Allocation Problems

- ▶ Locating a block can take many I/Os and disk seeks.
- ▶ Reliability can be a problem.
- ▶ The space required for the pointers.
  - Efficiency can be improved by clustering blocks into groups but increases internal fragmentation.

# File-Allocation Table (FAT)

- ▶ Beginning of volume has a table, indexed by block number.
- ▶ Much like a linked list, but faster on disk and cacheable.



# Indexed Allocation

# Indexed Allocation (1/2)

- ▶ **Indexed allocation:** each file has its own index block(s) of pointers to its data blocks.

# Indexed Allocation (1/2)

- ▶ **Indexed allocation**: each file has its own index block(s) of pointers to its data blocks.
- ▶ Need index table

# Indexed Allocation (1/2)


- ▶ **Indexed allocation**: each file has its own index block(s) of pointers to its data blocks.
- ▶ Need index table
- ▶ Random access



# Indexed Allocation (1/2)

- ▶ Indexed allocation: each file has its own index block(s) of pointers to its data blocks.
- ▶ Need index table
- ▶ Random access
- ▶ Dynamic access without external fragmentation, but have overhead of index block

## Indexed Allocation (2/2)



# Indexed Allocation Problems

- ▶ **Wasted space:** overhead of the index blocks.
- ▶ For example, even with a file of only one or two blocks, we need an entire index block.

- ▶ How large the index block should be?

# Index Block Size

- ▶ How **large** the **index block** should be?
- ▶ Keep the index block as **small** as possible.
  - We need a mechanism to hold pointers for **large files**.

# Index Block Size

- ▶ How **large** the **index block** should be?
- ▶ Keep the index block as **small** as possible.
  - We need a mechanism to hold pointers for large files.
- ▶ Mechanisms for this purpose include the following:
  - Linked scheme
  - Multi-level index
  - Combined scheme

- ▶ **Linked scheme:** link blocks of index table (no limit on size)

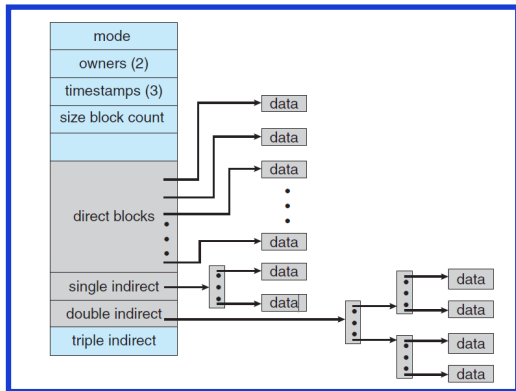
- ▶ Linked scheme: link blocks of index table (no limit on size)
- ▶ For example, an index block might contain a small header giving the name of the file and a set of the first 100 disk-block addresses.
- ▶ The next address is null or is a pointer to another index block.



- ▶ Two-level index
- ▶ A first-level index block to point to a set of second-level index blocks, which in turn point to the file blocks.
- ▶ Could be continued to a third or fourth level.

# Combined Scheme

- ▶ **Combine scheme**: used in Unix/Linux FS
- ▶ The first 12 pointers point to direct blocks
  - The data for small files do not need a separate index block.
- ▶ The next 3 pointers point to indirect blocks.
  - Single indirect
  - Double indirect
  - Triple indirect



- ▶ Best method depends on file access type.

- ▶ Best method depends on file access type.
- ▶ Contiguous is great for sequential and random.

- ▶ Best method depends on file access type.
- ▶ Contiguous is great for sequential and random.
- ▶ Linked is good for sequential, not random.

- ▶ Best method depends on file access type.
- ▶ **Contiguous** is great for **sequential and random**.
- ▶ **Linked** is good for **sequential, not random**.
- ▶ **Indexed** is more **complex**
  - Single block access could require 2 index block reads then data block read
  - Clustering can help improve throughput, reduce CPU overhead

# Summary

- ▶ FS layers: device, I/O control, basic FS, file-organization, logical FS, application



# Summary

- ▶ FS layers: device, I/O control, basic FS, file-organization, logical FS, application
- ▶ FS implementation:
  - On-disk structures: boot control block, volume control block, directory structure, and file control block
  - In-memory structures: mount table, directory structure, open-file tables, and buffers

# Summary

- ▶ FS layers: device, I/O control, basic FS, file-organization, logical FS, application
- ▶ FS implementation:
  - On-disk structures: boot control block, volume control block, directory structure, and file control block
  - In-memory structures: mount table, directory structure, open-file tables, and buffers
- ▶ Virtual file system (VFS)

# Summary

- ▶ FS layers: device, I/O control, basic FS, file-organization, logical FS, application
- ▶ FS implementation:
  - On-disk structures: boot control block, volume control block, directory structure, and file control block
  - In-memory structures: mount table, directory structure, open-file tables, and buffers
- ▶ Virtual file system (VFS)
- ▶ Directory implementation: linear list, and hash table

# Summary

- ▶ FS layers: device, I/O control, basic FS, file-organization, logical FS, application
- ▶ FS implementation:
  - On-disk structures: boot control block, volume control block, directory structure, and file control block
  - In-memory structures: mount table, directory structure, open-file tables, and buffers
- ▶ Virtual file system (VFS)
- ▶ Directory implementation: linear list, and hash table
- ▶ Allocation methods: contiguous allocation, linked allocation, and indexed allocation

# Questions?

## Acknowledgements

Some slides were derived from Avi Silberschatz slides.